

Appendix to "Enabling Almost Strong Consistency for Quorum-replicated Datastores"

Lingzhi Ouyang, Yu Huang, Hengfeng Wei, Jian Lu



1 Properties of Quorum-based Algorithms

In this section, we study the properties of quorum-based algorithms theoretically. For simplicity, we assume that all operations in the history are applied to the same register. Some conventions are specified as follows:

- \mathbb{W} : the set of write operations.
- \mathbb{R} : the set of read operations.
- $\mathbb{O} = \mathbb{W} \cup \mathbb{R}$: the set of operations including writes and reads.
- $[o_{st}, o_{ft}]$: the time interval of the operation o . The time of invocation event and response event of o are denoted o_{st} and o_{ft} separately.
- $ver(o)$: the version attached to the returned value if operation $o \in \mathbb{R}$, or the version attached to the written value if $o \in \mathbb{W}$.
- $ver(s, t)$: the version stored in the server s at time t .
- $maxVer(A, t) / minVer(A, t)$: the maximum / minimum version stored in the server set A at a specific time t .
- $maxVer(A) / minVer(A)$: the maximum / minimum version collected in a query, where A is the responding server set of the query. Note that each server in A may respond not at the same time.

The pseudo-code for client-server interaction appears in Algorithm 1. It can be easily proved that:

Lemma 1.1 (Query-after-Update Property). *For any update u and any query q , if $u \prec_{\sigma} q$, and q receives responses from the server set Q , then $maxVer(Q) \geq ver(u)$.*

Proof. According to the server procedures in Algorithm 1, the version stored in servers keep non-decreasing as time goes by. Therefore, for any update u , when it is finished at time t_1 , there exist a set of servers, denoted U , storing replicas with the version no lower than $ver(u)$. Namely, $minVer(U, t_1) \geq ver(u)$.

Now we consider any query q that starts after u . Namely, $u \prec_{\sigma} q$. Let Q denote the server set that q receives responses from.

According to the client procedures in Algorithm 1, any query or update will only be finished until a *majority* of servers reply, so $U \cap Q \neq \emptyset$. Assume the server $s \in U \cap Q$, and s returns the version at t_2 in the query q . Since q starts after u , then $t_2 > t_1$. Therefore, $maxVer(Q) \geq ver(s, t_2) \geq ver(s, t_1) \geq minVer(U, t_1) \geq ver(u)$. \square

The pseudo-codes for read/write operations appears in Algorithm 2 & 3. Here, we mainly focus on the properties in terms of monotonicity quorum-based algorithms, as shown in Table 1. These properties depict the relations between the temporal real-time relations and the semantic read-from relations of operations. The detailed proofs are mainly based on the *Query-after-Update Property*, and we display them in Section 1.1 - 1.4 respectively.

TABLE 1
Algorithm properties in terms of monotonicity

Properties	W2R2	W2R1	W1R2	W1R1
$w \prec_{\sigma} r \Rightarrow ver(w) \leq ver(r)$	✓	✓	✓	✓
$w \prec_{\sigma} w' \Rightarrow ver(w) < ver(w')$	✓	✓	×	×
$r \prec_{\sigma} r' \Rightarrow ver(r) \leq ver(r')$	✓	×	✓	×
$r \prec_{\sigma} w \Rightarrow ver(r) < ver(w)$	✓	×	×	×

1.1 Common Property: Write-Read Monotonicity

The quorum-based write / read algorithms with one or two communication round-trips guarantee the *write-read monotonicity* for multi-writer registers. Specifically,

Theorem 1.1 (Write-Read Monotonicity). *In quorum-based algorithms (see Algorithm 1, 2, 3), for any operations w, r on the same register in the execution history σ , where $w \in \mathbb{W}, r \in \mathbb{R}$, if $w \prec_{\sigma} r$, then $ver(w) \leq ver(r)$.*

Proof. $\forall w \in \mathbb{W}, r \in \mathbb{R}$, let u denote the update process of w , q denote the query process of r , and Q denote the server set that q receives responses from. Then $ver(w) = ver(u), ver(r) = maxVer(Q)$.

If $w \prec_{\sigma} r$, then, it's trivial to have $u \prec_{\sigma} q$. By Lemma 1.1, $maxVer(Q) \geq ver(u)$. Therefore, $ver(w) = ver(u) \leq maxVer(Q) = ver(r)$. \square

1.2 2-Round-Trip Write Property: Write-Write Monotonicity

The two-round-trip write algorithm guarantees the *write-write monotonicity* for multi-writer registers. Specifically,

Theorem 1.2 (Write-Write Monotonicity). *In quorum-based algorithms that completes each write operation in 2 communication round-trips (see TwoRoundWRITE in Algorithm 2), for any write operations w_1, w_2 on the same register in the execution history σ , if $w_1 \prec_{\sigma} w_2$, then $ver(w_1) < ver(w_2)$.*

Algorithm 1: Client-server interaction based on quorum communication.

```

1  $\triangleright$  Code for client process  $p_i (0 \leq i \leq n-1)$ :
2 function query( $key$ )
3    $vals \leftarrow \emptyset$ 
4   pfor each server  $s_j$   $\triangleright$  pfor: parallel for
5     send [ $'query', key$ ] to  $s_j$ 
6      $v \leftarrow [key, val, ver]$  from  $s_j$ 
7      $vals \leftarrow vals \cup v$ 
8   until a majority of them respond
9   return  $vals$ 
10 function update( $key, value, version$ )
11    $vals \leftarrow \emptyset$ 
12   pfor each server  $s_j$ 
13     send [ $'update', key, value, version$ ] to  $s_j$ 
14   wait for [ $'ACK'$ ]s from a majority of them
15  $\triangleright$  Code for server process  $s_i (0 \leq i \leq N-1)$ :
16 upon receive [ $'query', key$ ] from  $p_j$ 
17   send [ $'query - back', key, val, ver$ ] to  $p_j$ 
18 upon receive [ $'update', key, value, version$ ] from  $p_j$ 
19   pick  $[k, val, ver]$  with  $k == key$ 
20   if  $ver < version$  then
21      $val \leftarrow value$ 
22      $ver \leftarrow version$ 
23   send [ $'ACK'$ ] to  $p_j$ 

```

Algorithm 2: The write algorithms for client process $p_i (0 \leq i \leq n-1)$:

```

1 procedure TwoRoundWRITE( $key, value$ )
2    $replicas \leftarrow \text{query}(key)$ 
3    $version \leftarrow (\maxSeq(replicas) + 1, i)$ 
4   update( $key, value, version$ )
5 procedure OneRoundWRITE( $key, value$ )
6    $localSeq[key] \leftarrow localSeq[key] + 1$ 
7    $version \leftarrow (localSeq[key], i)$ 
8   update( $key, value, version$ )

```

Algorithm 3: The read algorithms for client process $p_i (0 \leq i \leq n-1)$:

```

1 procedure TwoRoundREAD( $key$ )
2    $replicas \leftarrow \text{query}(key)$ 
3    $version \leftarrow \maxVer(replicas)$ 
4    $value \leftarrow valWithMaxVer(replicas, version)$ 
5    $localSeq[key] \leftarrow version.seq$   $\triangleright$  Optional.
6   update( $key, value, version$ )
7   return  $value$ 
8 procedure OneRoundREAD( $key$ )
9    $replicas \leftarrow \text{query}(key)$ 
10   $version \leftarrow \maxVer(replicas)$ 
11   $value \leftarrow valWithMaxVer(replicas, version)$ 
12   $localSeq[key] \leftarrow version.seq$   $\triangleright$  Optional.
13  return  $value$ 

```

Proof. For two-round-trip write operations w_1, w_2 , let u_1 denote the update process of w_1 , q_2 denote the query process of w_2 , and Q_2 denote the server set that q_2 receives responses from. Then $ver(w_1) = ver(u_1), ver(w_2) > maxVer(Q_2)$ (Note: w_2 will construct a larger version according to the largest version in Q_2).

If $w_1 \prec_\sigma w_2$, then, it's trivial to have $u_1 \prec_\sigma q_2$. By Lemma 1.1, $maxVer(Q_2) \geq ver(u_1)$. Therefore, $ver(w_1) = ver(u_1) \leq maxVer(Q) < ver(w_2)$. \square

In comparison, the one-round-write algorithm can only guarantee the *write-write monotonicity* for *single-writer* registers, where the only writer can execute updates with *monotonically increasing local versions*. However, for *multi-writer* registers, one-round-write may result in *write-version inversion anomalies after a write*, which means that a *previous* write assigns a *larger* version for its written value than a *later* write (Namely, $(w_1 \prec_\sigma w_2) \cap (ver(w_1) > ver(w_2))$). What's worse, the procedure without the promise of assigning a version larger than previous writes may lead to unsuccessful updates on servers. More details will be discussed in Section 3 of this appendix.

1.3 2-Round-Trip Read Property: Read-Read Monotonicity

The two-round-trip read algorithm guarantees the *read-read monotonicity* for multi-writer registers. Specifically,

Theorem 1.3 (Read-Read Monotonicity). *In quorum-based algorithms that completes each read operation in 2 communication round-trips (see TwoRoundREAD in Algorithm 3), for any read operations r_1, r_2 on the same register in the execution history σ , if $r_1 \prec_\sigma r_2$, then $ver(r_1) \leq ver(r_2)$.*

Proof. For two-round-trip read operations r_1, r_2 , let u_1 denote the update process of r_1 , q_2 denote the query process of r_2 , and Q_2 denote the server set that q_2 receives responses from. Then $ver(r_1) = ver(u_1), ver(r_2) = maxVer(Q_2)$.

If $r_1 \prec_\sigma r_2$, then, it's trivial to have $u_1 \prec_\sigma q_2$. By Lemma 1.1, $maxVer(Q_2) \geq ver(u_1)$. Therefore, $ver(r_1) = ver(u_1) \leq maxVer(Q) = ver(r_2)$. \square

In comparison, the one-round-read algorithm can guarantee the *read-read monotonicity* for single-reader registers only if clients store previous read records locally. However, for multi-reader registers, the one-round-read algorithm may lead to *version inversion anomalies after a read*, which means that the version of a later operation (a read or a write) is smaller than the version of a *previous* read. The anomalies involves several patterns and more details will be discussed in Section 2 of this appendix.

1.4 Exclusive Property of W2R2: Read-Write Monotonicity

Besides all above properties, the W2R2 algorithm also guarantees the *read-write monotonicity* for MWMR registers.

Theorem 1.4 (Read-Write Monotonicity). *In the W2R2 algorithm, for any operations w, r on the same register in the execution history σ , where $w \in \mathbb{W}, r \in \mathbb{R}$, if $r \prec_\sigma w$, then $ver(r) < ver(w)$.*

Proof. $\forall w \in \mathbb{W}, r \in \mathbb{R}$, let u_r denote the update process of r , q_w denote the query process of w , and Q_w denote the server set that q receives responses from. Then $ver(r) = ver(u_r)$, $ver(w) > maxVer(Q_w)$ (Note: w will construct a larger version according to the largest version in Q_w).

If $w \prec_\sigma r$, then, it's trivial to have $u_w \prec_\sigma q_w$. By Lemma 1.1, $maxVer(Q_w) \geq ver(u_r)$. Therefore, $ver(r) = ver(u_r) \leq maxVer(Q_w) < ver(w)$. \square

2 Atomicity Violations of W2R1: Staleness and Probability

In this section, we first prove the possible atomicity violation patterns in W2R1 (see Algorithm 4), then we prove the bound of data staleness and calculate the probability of atomicity violations in W2R1 based on the violation patterns.

Algorithm 4: W2R1

```

1 Code for client process  $p_i (0 \leq i \leq n - 1)$ :
2 procedure TwoRoundWRITE( $key, value$ )
3    $replicas \leftarrow query(key)$ 
4    $version \leftarrow (maxSeq(replicas) + 1, i)$ 
5    $update(key, value, version)$ 
6 procedure OneRoundREAD( $key$ )
7    $replicas \leftarrow query(key)$ 
8    $version \leftarrow maxVer(replicas)$ 
9    $value \leftarrow valWithMaxVer(replicas, version)$ 
10  return  $value$ 

```

2.1 Proof of Atomicity Violation Patterns

We aim to prove that atomicity violations incurred in W2R1 are composed of RI or WI. When clients read and write a MWMR register using the W2R1 algorithm and obtain the history σ , we have that:

Theorem 2.1. *If σ violates atomicity, then there exists some operations in σ that form either RI or WI.*

Proof. If σ violates atomicity, then for any permutation π of σ , we have a stale read r , the dictating write w of r and the interfering write w' satisfying: $w \prec_\pi w' \prec_\pi r$. Then, we exhaustively check all cases that make r a stale read.

According to the definition of atomicity, two types of relations between operations are of our concern: the temporal real-time relation and the semantic read-from relation. We first enumerate all possible cases according to the semantic relation. Then for each case, we further enumerate all possible sub-cases according to the temporal relation.

Since w is the dictating write of r , we have that $ver(r) = ver(w)$. According to the semantic relation between versions of w and w' , we have two complementing cases: $ver(r) = ver(w) < ver(w')$ and $ver(r) = ver(w) > ver(w')$. In each case, we then consider the temporal relation between operations.

Case 1: $ver(r) = ver(w) < ver(w')$.

Note that r must be concurrent with w' . Namely, $r \parallel_\sigma w'$. This can be proved by contradiction. If $r \prec_\sigma w'$, w' can never be the interfering write of r . If $w' \prec_\sigma r$, according to the *write-read monotonicity* (Theorem 1.1), r must return the version of w' (or return an even larger version), contradicting the fact that $ver(r) = ver(w) < ver(w')$.

Next we exhaustively check all possible sub-cases according to the temporal relation to prove that: *there must exist a read operation r' that reads from w' , and r' precedes r* . Namely, $\exists r' = R(w') : r' \prec_\sigma r$.

We first enumerate all possible cases according to the temporal relation between w and w' . Since in the permutation π , $w \prec_\pi w'$, we have that in σ , $w \prec_\sigma w'$ or $w \parallel_\sigma w'$.

Case 1.1: $w \prec_\sigma w'$. The permutation between w and w' must be *determined* as $w \prec_\pi w'$. Then, we enumerate all cases as follows:

- × $\nexists r' = R(w')$. Note that $w' \parallel_\sigma r$, so r can be ordered before w' . Thus, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, which contradicts the permutation $w \prec_\pi w' \prec_\pi r$.
- × $\forall r' = R(w') : r' \not\prec_\sigma r$. Similarly, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, contradicting the permutation $w \prec_\pi w' \prec_\pi r$.
- ✓ $\exists r' = R(w') : r' \prec_\sigma r$. Then the permutation among w, w' and r must be $w \prec_\pi w'(\prec_\pi r') \prec_\pi r$ (see Fig. 1(a)).

Case 1.2: $w \parallel_\sigma w'$. In this case, the permutation between w and w' can not be *determined* only by the temporal relation of these two write operations.

Therefore, we then focus on other operations related to these two clusters¹. We first consider the temporal relation between w' and any read operation r'' that reads from w . Notice that r is a special instance of r'' , thus r'' must exist.

Similarly, by contradiction we have $w' \not\prec_\sigma r''$; otherwise r'' must return the version of w' (or return a even larger version) according to the *write-read monotonicity* (Theorem 1.1). Thus, the temporal relation between w' and r'' can only be $r'' \prec_\sigma w'$ or $r'' \parallel_\sigma w'$.

Case 1.2.1: $\exists r'' = R(w) : r'' \prec_\sigma w'$. By the temporal relation between r'' and w' , as well as read-from order between r'' and w , the permutation between w and w' must be *determined* as $w(\prec_\pi r'') \prec_\pi w'$.

- × $\nexists r' = R(w')$. Note that $w' \parallel_\sigma r$, so r can be ordered before w' . Thus, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, contradicting the permutation $w \prec_\pi w' \prec_\pi r$.
- × $\forall r' = R(w') : r' \not\prec_\sigma r$. Similarly, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, contradicting the permutation $w \prec_\pi w' \prec_\pi r$.
- ✓ $\exists r' = R(w') : r' \prec_\sigma r$. Then the permutation must be $w \prec_\pi w'(\prec_\pi r') \prec_\pi r$ (see Fig. 1(?)).

Case 1.2.2: $\forall r'' = R(w) : r'' \parallel_\sigma w'$. Then, We consider the temporal relation between w and any read operation r' that reads from w' . Note that r' may not exist.

Case 1.2.2.1: $\nexists r' = R(w')$. Then, there exists no determining factor to deciding the permutation between w and w' . Thus, the permutation among w, w', r is $w' \prec_\pi w \prec_\pi r$ or $w \prec_\pi r \prec_\pi w'$, both contradict the assumption that w, w' and r form an inconsistent read. [×]

Case 1.2.2.2: $\exists r' = R(w') : r' \prec_\sigma w$. By the temporal relation between r' and w , as well as read-from order between r' and w' , the permutation between w and w' is *determined* as $w'(\prec_\pi r') \prec_\pi w$, which contradicts $w \prec_\pi w'$. [×]

1. Here, we use the concept *cluster*, a terminology introduced by Gibbons and Korach [1], for our analysis. A *cluster* is a subset of operations in an execution that consists of a write and all of its dictated reads. Gibbons and Korach [1] have shown that atomicity violations indicates certain rules related to different zones of clusters. Therefore, for soundness of proof, atomicity violations concerned with w, w', r should be analyzed with all operations of their related clusters, consisting of w and all of its dictated reads, as well as w' and all of its dictated reads, because any dictating reads that returns related written values may effect the permutation.

Case 1.2.2.3: $\forall r' = R(w') : r' \not\prec_\sigma w$. Then, there exists no determining factor to deciding the permutation between w and w' .

- × $\forall r' = R(w') : r' \not\prec_\sigma r$. Thus, the permutation among w, w' and r is $w' \prec_\pi w \prec_\pi r$ or $w \prec_\pi r \prec_\pi w'$, both violating the given permutation $w \prec_\pi w' \prec_\pi r$.
- ✓ $\exists r' = R(w') : r' \prec_\sigma r$. Then the permutation can be $w \prec_\pi w' \prec_\pi r' \prec_\pi r$ or $w' \prec_\pi w \prec_\pi r' \prec_\pi r$ (see Fig. 1(b)), which involves atomicity violations anyway.

Above all, for $ver(w) < ver(w')$, then w, w' and r can be linearly extended to $w \prec_\pi w' \prec_\pi r$ only when $\exists r' = R(w') : r' \prec_\sigma r$. Obviously, r and r' form RI.

Case 2: $ver(r) = ver(w) > ver(w')$.

Note that w must be concurrent with w' . This can be proved by contradiction. If $w \prec_\sigma w'$, according to *write-write monotonicity* that two round-trip write algorithm holds (Theorem 1.2), the version of w' must be larger. If $w' \prec_\sigma w$, we can never get $w \prec_\pi w' \prec_\pi r$ when linearly extending σ to π . Besides, according to the permutation $w' \prec_\pi r$, we have $w' \prec_\sigma r$ or $w' \parallel_\sigma r$.

Next we enumerate all possible sub-cases according to the temporal relation to prove that: there exists a read operation r'' (can be r) that reads from w , satisfying that r'' precedes either w' or another read operation r' that reads from w' . Namely, $\exists r'' = R(w) : (r'' \prec_\sigma w') \vee (\exists r' = R(w') : r'' \prec_\sigma r')$.

We first consider the temporal relation between w' and any read operation r'' that reads from w . Notice that r is a special instance of r'' , thus r'' must exist.

Case 2.1: $\exists r'' = R(w) : r'' \prec_\sigma w'$. Then, the permutation between w and w' must be *determined* as $w(\prec_\pi r'') \prec_\pi w'$. Note that r cannot take the role of r'' here, because $r \not\prec_\sigma w'$.

Case 2.1.1: $w' \prec_\sigma r$. Then the permutation among w, w' and r must be $w \prec_\pi w' \prec_\pi r$ (see Fig. 1(?)). [✓]

Case 2.1.2: $w' \parallel_\sigma r$.

- × $\nexists r' = R(w')$. Note that $w' \parallel_\sigma r$, so r can be ordered before w' . Thus, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, contradicting the assumption.
- × $\forall r' = R(w') : r' \not\prec_\sigma r$. Similarly, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, contradicting the assumption.
- ✓ $\exists r' = R(w') : r' \prec_\sigma r$. Then the permutation among w, w' and r must be $w(\prec_\pi r'') \prec_\pi w'(\prec_\pi r') \prec_\pi r$ (see Fig. 1(?)).

Case 2.2: $\forall r'' = R(w) : r'' \not\prec_\sigma w'$. Then, We consider the temporal relation between r'' and any read operation r' that reads from w' .

Case 2.2.1: $\nexists r' = R(w')$. Then, there exists no determining factor to deciding the permutation between w and w' .

- × $w' \prec_\sigma r$. Then the permutation among w, w' and r must be $w' \prec_\pi w \prec_\pi r$, contradicting the assumption.
- × $w' \parallel_\sigma r$. Thus, the permutation among w, w', r is $w' \prec_\pi w \prec_\pi r$ or $w \prec_\pi r \prec_\pi w'$, both contradicting the assumption.

Case 2.2.2: $\forall r' = R(w'), r'' = R(w) : r'' \not\prec_\sigma r'$.

- × $w' \prec_\sigma r$. Then the permutation among w, w' and r must be $w' \prec_\pi w \prec_\pi r$, contradicting the assumption.

$\times w' \parallel_{\sigma} r$. Thus, the permutation among w, w', r is $w' \prec_{\pi} w \prec_{\pi} r$ or $w \prec_{\pi} r \prec_{\pi} w'$, both contradicting the assumption.

Case 2.2.3: $\exists r' = R(w'), r'' = R(w) : r'' \prec_{\sigma} r'$.

By contradiction we have $w \not\prec_{\sigma} r'$; otherwise r' must return the version of w (or return a even larger version) according to the *write-read monotonicity* (Theorem 1.1). Besides, r'' can't precede its dictating write w due to the regularity property. Thus, the temporal relation can only be $w \parallel_{\sigma} r''$ as well as $w \parallel_{\sigma} r'$. The permutation among w, r'', r' must be $w \prec_{\pi} r'' \prec_{\pi} r'$. Then the permutation can be $w \prec_{\pi} w' \prec_{\pi} r'' \prec_{\pi} r'$ or $w' \prec_{\pi} w \prec_{\pi} r'' \prec_{\pi} r'$ (see Fig. 1(d)), which involves atomicity violations anyway. If $w' \prec_{\sigma} r$ (or r just takes the role of r''), then r may become a stale read (see Fig. 1(?)). [✓]

Above all, for $ver(w) > ver(w')$, then w, w' and r can be linearly extended to $w \prec_{\pi} w' \prec_{\pi} r$ only when there exists a read operation r'' (can be r) that reads from w , satisfying that r'' precedes either w' or another read operation r' that reads from w' . Namely, $\exists r'' = R(w) : (r'' \prec_{\sigma} w') \vee (\exists r' = R(w') : r'' \prec_{\sigma} r')$, where r'' and w' form WI, or r'' and r' form RI.

In conclusion, by assuming r is a stale read, we exhaustively check all cases according to the temporal real-time relations and the semantic read-from relations of relative operations. At last, we have proved that the necessary conditions to make r a stale read involve RI or WI. Therefore, Theorem 2.1 is true. \square

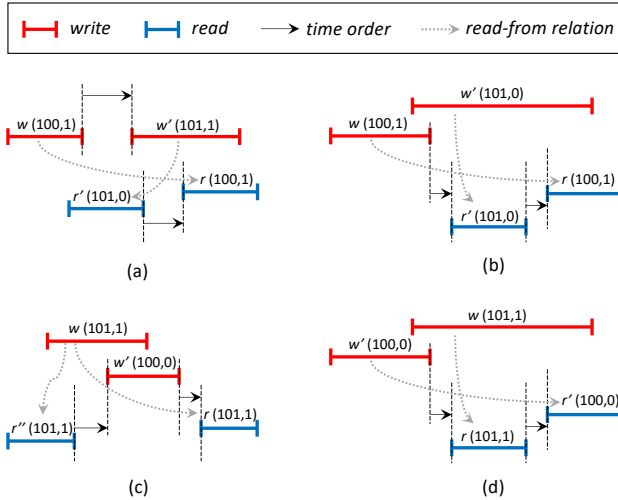


Fig. 1. Examples of atomicity violations that make $w \prec_{\pi} w' \prec_{\pi} r, r = R(w)$ in W2R1. Note: (a) & (b) present the possible cases when $ver(w) < ver(w')$, while (c) & (d) present the possible cases when $ver(w) > ver(w')$. The traces in (a) & (c) are bound to produce the permutation involving $w \prec_{\pi} w' \prec_{\pi} r$, while the traces in (b) & (d) can produce other optimal permutations besides those involving $w \prec_{\pi} w' \prec_{\pi} r$. However, for the Theorem integrity, we should take these possible cases into account in our proof.

From the proof of Theorem 2.1 we have a more detailed proposition, which is useful for our further analysis of staleness bound.

Specifically, then for any permutation π of σ , we have a stale read r , the dictating write w of r and the interfering write w' satisfying: $w \prec_{\pi} w' \prec_{\pi} r$.

Proposition 2.1. Suppose σ is the history obtained when clients read and write a MWMM register using the W2R1 algorithm. For any permutation π of σ , there exists a stale read r , the dictating write w of r and the interfering write w' satisfying: $w \prec_{\pi} w' \prec_{\pi} r$, as well as

- **Case 1:** $ver(w) < ver(w')$. w' and r are concurrent, and there exists a read r' that reads from w' , as well as precedes r . Namely, $(w' \parallel_{\sigma} r) \wedge (\exists r' = R(w') : r' \prec_{\sigma} r)$;
- **Case 2:** $ver(w) > ver(w')$. w and w' are concurrent, and there exists a read r'' (can be r) that reads from w , satisfying that r'' precedes either w' or another read r' that reads from w' . Namely, $(w \parallel_{\sigma} w') \wedge (\exists r'' = R(w) : (r'' \prec_{\sigma} w') \vee (\exists r' = R(w') : r'' \prec_{\sigma} r'))$.

2.2 Proof of Bound of Data Staleness

In this section, we calculate the tight bound of data staleness when accessing a MWMM register using the W2R1 algorithm. Denote the number of writer clients as n_w . For any history σ , we have that:

Theorem 2.2. There exists a linear extension π of σ such that any read in π returns the value of the latest B preceding write. Here, $B = n_w + \frac{1}{2}n_w(n_w - 1) + 1$. Moreover, the bound B is tight, i.e. there exists a permutation π in which some read returns values of the oldest write in the latest B preceding writes.

Proof. Suppose the distributed storage system consists of N replicas ($N \geq 3$). We divide n clients into n_r readers and n_w writers, meaning that each client can raise either read or write requests, but not both. Such simplification won't have essential effects on performance proof or probabilistic analysis². These n clients are allowed to apply W2R1 to access the distributed storage system and an execution σ will be produced over time.

Suppose r is a read in σ that returns a *stale* value written by w . According to theorem ??, there must exist another write operation, denoted w' , that makes r return a *stale* value. Formally, $\exists w' : w \prec_{\pi} w' \prec_{\pi} r$ where $r = R(w)$. Then, the proof of upper bound of k in terms of k -atomicity can be transferred into the following problem: given a read r and its dictating write w in σ with n_w writer processes, how many different writes w' at most that satisfy $w \prec_{\pi} w' \prec_{\pi} r$? We solve it by following the case-by-case proof of theorem 2.1. Through exhausting all possible cases according to *version relations* among w, w' and r separately, the upper bound of k is to consider all worst cases together.

Case 1: $ver(w) = ver(r) < ver(w')$. From P1 we know w' and r are concurrent, and there exists a read r' that reads from w' , as well as precedes another read r'' that reads from w . In this pattern, w' is the dictating write for another read r' that happens before r , which indicates the invocation of w' is *earlier* than that of r . Consider the details from the perspective of quorum-replicated communications. $r = R(w)$ indicates that *at some point* before the read operation r is done, there exists a majority of replicas storing the value with a version no larger than $ver(w)$. Thus, candidate writes that can act the role of w' are definitely stay unfinished *at*

2. Actually, the simplification will be even more adverse to the algorithm compared to the model that each client can initiate both read and write requests.

that point of time. On the one hand, each writer are able to invoke at most one write containing *that special point*; on the other hand, all writes that starts after *that special point* can't dictate any read previous to r , so these writes and dictating reads can be mapped *after* r in the permutation safely. Therefore, the number of candidate writes is no more than the number of writers n_w . Meanwhile, each of the candidates has some dictating read(s) previous to r . Fig.2 shows the overall construction for this case.




Fig. 2. The construction of worst case of atomicity violations in W2R1 for $ver(w) < ver(w')$.

Case 2: $ver(w) = ver(r) > ver(w')$. P2 provides two patterns that are allowed, and we use the following one: w and w' are concurrent, and there exists a read r'' that reads from w and precedes w' . This pattern is selected because it can help us to construct the worst case of atomicity violations. In order to find out how many different writes at most are legal to act the role of w' whose version is smaller than $ver(w)$, we use *adversary argument* to fill as many writes after r'' as possible. Suppose the read r returns a value versioned (seq, i) . According to Algorithm 2, the two-round-trip dictating write will assign the version (seq, i) unless there exists some replica storing a value versioned $(seq - 1, id)$, where id can be any writer's identifier. Similarly, the version $(seq - 1, id)$ won't be assigned unless there exists a value versioned $(seq - 2, id')$ in some replica, and so on. Then, the adversary argument is that each writer assigns a version whose sequence number is increased by another version from a concurrent write. We call these writes *the steps to* (seq, i) . When r'' accesses replicas, there exists at least a majority of replicas where (seq, i) is the maximum version attached to the value. After r'' is finished, those stepping-writes are then completed in an adversary sequence. Each writer has at most one above stepping-write and may have subsequent writes whose version is larger than that of its stepping-write and smaller than (seq, i) , and all these subsequent writes are legal to act the role of w' .³ From the construction of the worst case, we get candidates for w in the case of $ver(w') < ver(w)$ is no more than $\frac{1}{2}n_w(n_w - 1)$.⁴ Fig.3 shows the overall construction for $ver(w) > ver(w')$.

Above all, the upper bound of k is to consider two worst cases together. Fig.4 in Appendix A shows an example of the




Fig. 3. The construction of worst case of atomicity violations in W2R1 for $ver(w) > ver(w')$. Note: the subscript of write in the shaded area only denotes the writer identifier. The steps are able to be ordered before w in π . All writes denoted w' are not only after r'' but also before r . Since both r'' and r read from w , so all writes denoted w' can only be ordered after w and r'' as well as before r in π .

worst case in W2R1. The upper bound of k is proved to be

$$\begin{aligned} k_{max} &= |\{w' | ver(w') < ver(w)\}|_{max} + \\ &\quad |\{w' | ver(w') > ver(w)\}|_{max} + 1 \\ &= \frac{1}{2}n_w(n_w - 1) + n_w + 1 \\ &= \frac{1}{2}n_w(n_w + 1) + 1, \end{aligned}$$

which can only be achieved by the client process with the *maximum identifier*.

When $N < 3, \dots$

Here we present an example of worst case. Suppose the distributed storage system contains 3 replicas. Clients include 3 writers and 3 readers. Suppose all operations are performed on a single register. One example of the worst case in W2R1 displays in Fig.4. The details on replicas are described in Table 2.

Till now, we prove Theorem 2.2 that W2R1 algorithm achieves the emulation of k -atomic, multi-writer multi-reader registers, where k is bounded by the number of writers n_w . Specially, W2R1 satisfies 2-atomicity for *single-writer multiple-reader* registers, which conforms to the conclusion in [2]. \square

3. Note that when we try to construct the worse case of atomicity violations for r , we should think in adversary argument. In this case, that is, try to shrink the number of w' as much as possible. Therefore, those stepping-writes should be ordered before r'' in π even if they have some dictated reads happening after r'' .

4. It can be proved that *the steps to* (seq, i) formed by all writers can enable most write candidates that are legal to act the role of w' in the case of $ver(w') < ver(w)$, even compared with other patterns.

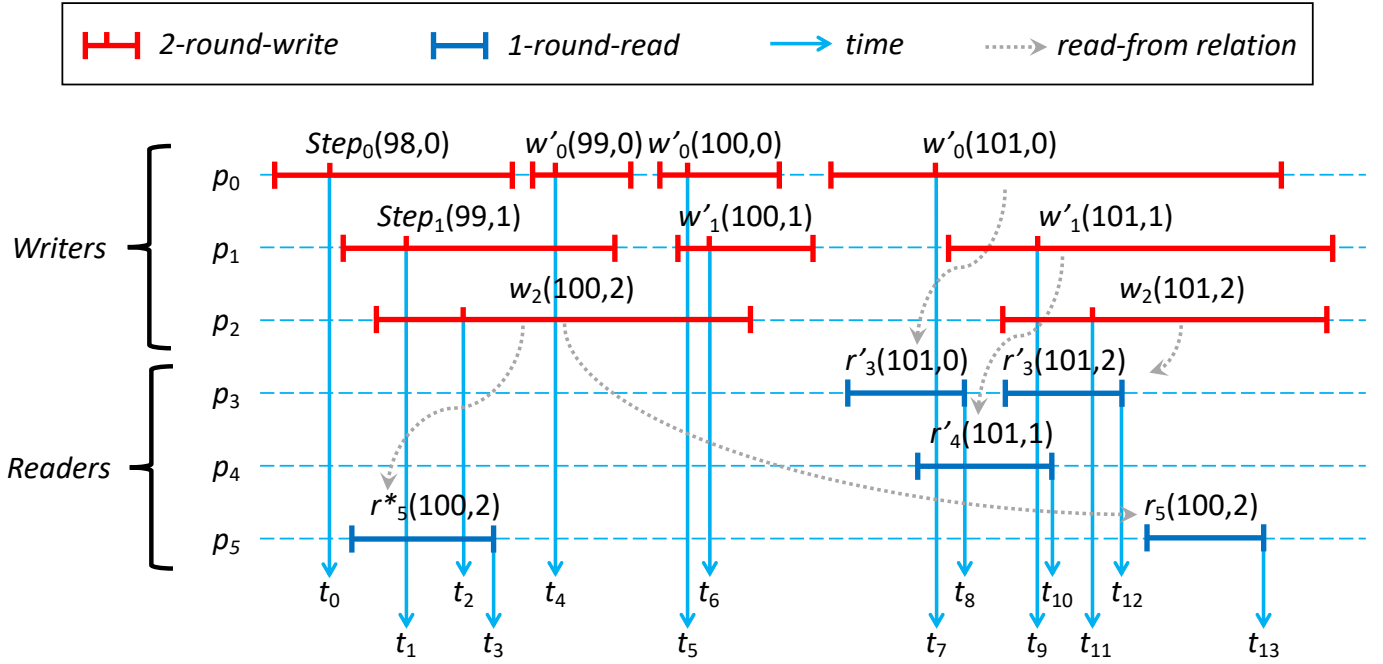


Fig. 4. An example of worst case in W2R1. The version that a write assigns or a read returns are displayed in brackets. One permutation of these operations can be: $Step_0(98,0) \prec_{\pi} Step_1(99,1) \prec_{\pi} w_2(100,2) \prec_{\pi} r_5''(100,2) \prec_{\pi} w'_0(99,0) \prec_{\pi} w'_0(100,0) \prec_{\pi} w'_1(100,1) \prec_{\pi} w'_0(101,0) \prec_{\pi} r'_3(101,0) \prec_{\pi} w'_1(101,1) \prec_{\pi} r'_3(101,2) \prec_{\pi} r_5(100,2)$. For $r_5(100,2)$, it returns the value written by one of the latest 7 preceding writes in π on the same key. Legal writes for acting the role of w' include $w'_0(99,0), w'_0(100,0), w'_1(100,1), w'_0(101,0), w'_1(101,1), w'_2(101,2)$. Among these writes, $w'_0(99,0), w'_0(100,0), w'_1(100,1)$ together with $w_2(100,2), r_5(100,2)$ and $r_5''(100,2)$ satisfy the form of Pattern (2). $w'_0(101,0), w'_1(101,1), w'_2(101,2)$ with their dictating reads as well as $w_2(100,2)$ and $r_5(100,2)$ satisfy the form of Pattern (5).

TABLE 2
Details on Replicas inn Fig.4

Version	Replica 1	Replica 2	Replica 3	Description
t_0	(97,0)	(97,0)	(97,0)	p_0 then assigns (98,0) after the query returns maximum version (97,0) in 1 st round-trip.
t_1	(98,0)	(97,0)	(97,0)	p_1 then assigns (99,1) after the query returns maximum version (98,0) in 1 st round-trip.
t_2	(99,1)	(97,0)	(97,0)	p_2 then assigns (100,2) after the query returns maximum version (99,1) in 1 st round-trip.
t_3	(100,2)	(97,0)	(97,0)	p_5 returns the value versioned (100,2) in the read.
t_4	(100,2)	(98,0)	(97,0)	p_0 then assigns (99,0) after the query returns maximum version (98,0) in 1 st round-trip.
t_5	(100,2)	(99,0)	(97,0)	p_0 then assigns (100,0) after the query returns maximum version (99,0) in 1 st round-trip.
t_6	(100,2)	(99,0)	(97,0)	p_1 then assigns (100,1) after the query returns maximum version (99,1) in 1 st round-trip.
t_7	(100,2)	(100,2)	(97,0)	p_0 assigns (101,0) after the query returns maximum version (100,2) in 1 st round-trip.
t_8	(100,2)	(101,0)	(97,0)	p_3 returns the value versioned (101,0) in the read.
t_9	(100,2)	(101,0)	(97,0)	p_1 assigns (101,1) after the query returns maximum version (100,2) in 1 st round-trip.
t_{10}	(100,2)	(101,1)	(97,0)	p_4 returns the value versioned (101,1) in the read.
t_{11}	(100,2)	(101,1)	(97,0)	p_2 assigns (101,2) after the query returns maximum version (100,2) in 1 st round-trip.
t_{12}	(100,2)	(101,2)	(97,0)	p_3 returns the value versioned (101,2) in the read.
t_{13}	(100,2)	(101,2)	(97,0)	p_5 returns the value versioned (100,2) in the read.

2.3 Probability of Atomicity Violations

In this subsection we quantify the atomicity violations incurred in W2R1. The quantification follows from theorem 2.1 that, atomicity violations incurred in W2R1 are composed of either RI or WI. We simplify the quantification problem into the following one: what's the probability that RI or WI incurred in W2R1? The simplification helps us to quantify atomicity violations by ignoring unnecessary details. Note that the occurrence of either RI or WI is merely the condition *necessary* but *not sufficient* for atomicity violations in W2R1(Fig. ??). Thus, the probabilistic results of atomicity violations by quantifying the occurrence of both patterns

would be larger than the exact model. That is,

$$\begin{aligned}
 \mathbb{P}\{\text{Atomicity Violations in W2R1}\} &\leq \mathbb{P}\{\text{RI} \vee \text{WI}\} \\
 &= \mathbb{P}\{\text{RI}\} + \mathbb{P}\{\text{WI}\} - \mathbb{P}\{\text{RI} \wedge \text{WI}\} \\
 &\leq \mathbb{P}\{\text{RI}\} + \mathbb{P}\{\text{WI}\}
 \end{aligned} \tag{2.1}$$



Fig. 5. The relation between atomicity violations and two types of version inversions in W2R1 .

According to the case-by-case proof of theorem 2.1, RI and WI can only occur within certain conditions in W2R1. Intuitively, an execution with higher degree of concurrency would produce more atomicity violations, but it's still not sufficient without certain read-write patterns. Thus, we decompose RI and WI incurred in W2R1 into a *concurrency pattern* and a *read-write pattern* separately. The *concurrency pattern*(CP) describes the constraints of real-time orders among events of operations, and the *read-write pattern*(RWP) limits read/write semantics among operations. Then, we have

Definition 2.1. *The read version inversion after a read in W2R1 (RI-W2R1) are composed of one write w' and two reads r, r' , satisfying the requirements of*

- *the long-lived-write concurrency pattern*(CP):
 - 1) $r_{st} \in [w'_{st}, w'_{ft}]$,
 - 2) $r'_{ft} \in [w'_{st}, r_{st}]$;
- *the non-monotonic read-write pattern*(RWP):
 - 3) $r' = R(w')$;
 - 4) $r \neq R(w')$.

Definition 2.2. *The write-version inversion after a read in W2R1 (WI-W2R1) are composed of two writes w, w' and one read r'' , satisfying the requirements of*

- *the long-lived-write concurrency pattern*(CP):
 - 1) $w'_{st} \in [w_{st}, w_{ft}]$,
 - 2) $r''_{ft} \in [w_{st}, w'_{st}]$;
- *the non-monotonic read-write pattern*(RWP):
 - 3) $r'' = R(w)$;
 - 4) $ver(w') \neq ver(w)$.

For convenience, we use RI and WI to refer to RI-W2R1 and WI-W2R1 in this subsection. Exactly speaking, the conditions described in Definition 2.1 & 2.2 are *necessary* but *not sufficient* for RI or WI since they ignore the version relations among operations. However, these conditions are mean enough for our evaluation.

Here we use Wei's approach [2] to calculate the rate of CP and RWP. In [2], CP and RWP are quantified with a *queueing model* and a *timed balls-into-bins model* respectively. These two models have been proposed to quantify *old-new inversion anomalies*(ONI) in *single-writer multi-reader*(SWMR) registers. From [2] we have

$$\begin{aligned}
 \mathbb{P}\{\text{ONI}\} &= \sum_{m \geq 1} \mathbb{P}\{\text{CP} \mid R' = m\} \cdot \mathbb{P}\{\text{RWP} \mid R' = m\} \\
 &\approx \sum_{m=1}^{n_r} \left(\left(\sum_{k=0}^{n_r-1} \binom{n_r}{k} \binom{m-1}{n_r-k-1} p_0^{k n_r-k} s^m \right) \right. \\
 &\quad \cdot e^{-q\lambda_w t} \frac{\alpha^q B(q, \alpha(N-q)+1)}{B(q, N-q+1)} \\
 &\quad \cdot \left(1 - \left(\frac{J_1}{B(q, N-q+1)} \right)^m \right) \Bigg). \quad (2.2)
 \end{aligned}$$

where R' is a random variable denoting the number of r 's in Definition 2.1. It is an approximation since the timed balls-into-bins model used for calculating the probability of read-write patterns assumes that there is at most one such r' in each reader queue, which can partly be justified by numerical results.

Through observation we can see that ONI is a special case of RI when there exists single writer. Based on the quantification of ONI, we can quantify RI for MWMR registers with extra consideration of the concurrency effects of multiple writers. As for WI, by replacing r, w', r' in Definition 2.1 with w', w, r'' accordingly, it can be analyzed with similar process. Above all, we simplify MWMR quantification by reusing results of ONI and taking the extra effects of multiple writers into consideration together. Such simplification would lead to a more relaxed evaluation results, because it calculates the effects of r' repeatedly. However, since it's infeasible to consider all possible combinations of real-time orders and read-write semantics among operations, the simplification is necessary.

The extra effects of multiple concurrent writers are described with the following queueing model. Each client's workload is recognized as an independent queue characterized by the rate of operations and the seRice time of each operation (i.e. $[o_{st}, o_{ft}]$). Assume a Poisson process with parameter λ for the scenario of each client issuing a sequence of read/write operations, and an exponential distribution with parameter μ for the seRice time of each operation. We then have n independent, parallel $M/M/1/1/\infty/FCFS$ queues (i.e., a single-server exponential queueing system, whose capacity is 1 with the "first come first served" discipline), all with arrival rate λ and seRice rate μ . For simplicity, queues of writes and reads are separate; and if there is any operation in seRice, no more operations can enter it in that queue.

Let $X^i(t)$ be the number of operations in queue i at time t . Then $X^i(t)$ is a continuous-time Markov chain with two states: 0 when the queue is empty and 1 when some operation is being served. Its stationary distribution $P_s \triangleq P(X^i(\infty) = s)$, $s \in \{0, 1\}$ is:

$$P_0 = \frac{\mu}{\mu + \lambda}, P_1 = \frac{\lambda}{\mu + \lambda}.$$

Now we consider the concurrency conditions of RI and WI respectively. For RI, given a read r in Q_i , let W'_{cr} be a random variable denoting the number of w 's satisfying $r_{st} \in [w'_{st}, w'_{ft}]$ in RI. The probability of the event that r starts during the seRice period of some write w' in Q_j equals the probability that when r arrives Q_i , it finds Q_i empty, as well as finds Q_j full as a bystander (with the constraint that Q_j is a writer queue). Since the events in different queues are independent, by the PASTA property and through combinatorial analysis, we have

$$\mathbb{P}\{W'_{cr} = w\} = \binom{n_w}{w} \left(\frac{\lambda}{\mu + \lambda} \right)^w \left(\frac{\mu}{\mu + \lambda} \right)^{n_w - w + 1} \quad (2.3)$$

Conditioning on $W'_{cr} = w$, the probability of RI is no more than the sum of each w' forming RI with r separately. Therefore, a probabilistic and combinatorial analysis shows that

$$\begin{aligned}
 \mathbb{P}\{\text{RI}\} &= \sum_{w=1}^{n_w} \mathbb{P}\{W'_{cr} = w\} \times \mathbb{P}\{\text{RI} \mid W'_{cr} = w\} \\
 &\leq \sum_{w=1}^{n_w} \mathbb{P}\{W'_{cr} = w\} \times w \mathbb{P}\{\text{ONI}\} \quad (2.4)
 \end{aligned}$$

As for WI, given a specific write w' in Definition 2.2, let $W_{cw'}$ be a random variable denoting the number of writes w

satisfying $w'_{st} \in [w_{st}, w_{ft}]$ in RI. Similarly, through probabilistic and combinatorial analysis we have

$$\mathbb{P}\{W_{cw'} = w\} = \binom{n_w - 1}{w} \left(\frac{\lambda}{\mu + \lambda}\right)^w \left(\frac{\mu}{\mu + \lambda}\right)^{n_w - w} \quad (2.5)$$

$$\begin{aligned} \mathbb{P}\{\text{WI}\} &= \sum_{w=1}^{n_w-1} (\mathbb{P}\{W_{cw'} = w\} \times \mathbb{P}\{\text{ONI} \mid W_{cw'} = w\}) \\ &\leq \sum_{w=1}^{n_w-1} (\mathbb{P}\{W_{cw'} = w\} \times w \mathbb{P}\{\text{ONI}\}) \end{aligned} \quad (2.6)$$

Substituting Formula (2.2)-(2.6) into (2.1), we obtain an upper bound of the rate of violating atomicity incurred in W2R1.

$$\begin{aligned} &\mathbb{P}\{\text{Atomicity Violations in W2R1}\} \\ &\leq \sum_{x=1}^{n_w} \mathbb{P}\{W'_{cr} = x\} \cdot x \mathbb{P}\{\text{ONI}\} + \sum_{x=1}^{n_w-1} \mathbb{P}\{W_{cw'} = x\} \cdot x \mathbb{P}\{\text{ONI}\} \\ &\approx \frac{(2n_w - 1)\lambda\mu}{(\lambda + \mu)^2} \cdot \sum_{m=1}^{n_r} \left(\left(\sum_{k=0}^{n_r-1} \binom{n_r}{k} \binom{m-1}{n_r-k-1} p_0^k r^{n_r-k} s^m \right) \right. \\ &\quad \cdot e^{-q\lambda_w t} \frac{\alpha^q B(q, \alpha(N-q+1))}{B(q, N-q+1)} \cdot \left(1 - \left(\frac{J_1}{B(q, N-q+1)} \right)^m \right) \end{aligned}$$

Note again that the occurrence of RI or WI is not enough to result in atomicity violations, so the probability would be larger than that of atomicity violations.

From the numerical analysis, in which we have chosen $\lambda = \mu = 10s^{-1}$ and $\lambda_r = \lambda_w = 20s^{-1}$ based on the experimental results in Section ??, we observe that (see Fig. ...; the detailed results can be found in supplementary Appendix 4):

Observation 2.1. *Probabilistically, the W2R1 algorithm rarely violates atomicity.*

Observation 2.2. *The rate of atomicity violations incurred in the W2R1 algorithm has positive correlation to the number of concurrent writers.*

3 W1R2 & W1R1

From the diamond structure we have one-round-write algorithms W1R2 & W1R1 (Algorithm 5 & 6), which complete each write in only one communication round-trip. As is shown in Algorithm ??, local storage is required for each writer process in one-round-write. When a writer initiates a write request, it would construct a version for the written key-value pair based on its local sequence number. The version assigned for its written value is only guaranteed to be larger than any version returned by its previous reads or assigned by its previous writes, but not others. Therefore, one-round-write works well only for *single-writer* registers since the only writer can promise to preserve the version for a register monotonically increasing without conflicting with other writers. For *multiple-writer* registers, one-round-write may result in *write-version inversion anomalies*: a later write updates a value with a lower version than a previous write from a different writer. Here, we define a write is *invisible* if the assigned version of the write is lower than a quorum of replicas before the invocation of the write. The value updated by an invisible write cannot be return by any read.

In this section, we tend to demonstrate that one-round-write algorithms are useless from two aspects. First we prove both W1R2 and W1R1 don't guarantee the value returned by each read can be within bounded staleness in Section 3.1. Then we argue that both algorithms lead to frequent unsuccessful updates by quantifying the rate of invisible writes in Section 3.2.

Algorithm 5: W1R2 without acknowledgement

```

1 Code for client process  $p_i (0 \leq i \leq n-1)$ :
2 procedure OneRoundWRITE( $key, value$ )
3    $localSeq[key] \leftarrow localSeq[key] + 1$ 
4    $version \leftarrow (localSeq[key], i)$ 
5   update( $key, value, version$ )
6 procedure TwoRoundREAD( $key$ )
7    $replicas \leftarrow query(key)$ 
8    $version \leftarrow \maxVer(replicas)$ 
9    $value \leftarrow valWithMaxVer(replicas, version)$ 
10   $localSeq[key] \leftarrow version.seq$ 
11  update( $key, value, version$ )
12  return  $value$ 
```

Algorithm 6: W1R1 without acknowledgement

```

1 procedure OneRoundWRITE( $key, value$ )
2    $localSeq[key] \leftarrow localSeq[key] + 1$ 
3    $version \leftarrow (localSeq[key], i)$ 
4   update( $key, value, version$ )
5 procedure OneRoundREAD( $key$ )
6    $replicas \leftarrow query(key)$ 
7    $version \leftarrow \maxVer(replicas)$ 
8    $value \leftarrow valWithMaxVer(replicas, version)$ 
9    $localSeq[key] \leftarrow version.seq$ 
10  return  $value$ 
```

3.1 Proof of Atomicity Violations in One-Round-Write

[证明：一轮写算法不保证有界陈旧度]

We prove that, in the W1R2 or W1R1 algorithm, the value returned by each read can be any stale. That is,

Theorem 3.1. *For MWMM registers, given $k \in \mathbb{N}^*$, the W1R2 or W1R1 algorithm doesn't satisfy k -atomicity.*

Proof. Assume W1R2 or W1R1 satisfies k -atomicity, where $k \in \mathbb{N}^*$. Then, we can raise a counterexample to show the absurdity of above assumption. The counterexample is constructed as follows. Suppose there exists two writers whose identifiers are p_0 and p_1 , and we only focus on the same register. At time t_0 , p_0 stores v_0 as its local sequence number for the register, and p_1 stores v_1 , satisfying $v_1 > v_0$. During the period from t_0 to t_1 , p_1 executes k writes successively while p_0 initiates no writes. Thus, at time t_1 , p_1 's local version is increased to be $(v_1 + k, 1)$ while p_0 still keeps v_0 . After that, during the period from t_1 to t_2 , only p_0 executes k writes successively and its local version is increased to be $v_0 + k$ at t_2 . Since $(v_0 + k, p_0) < (v_1 + k, p_1)$, the written value versioned $(v_0 + k, p_0)$ is invisible. Then at t_3 , a read occurs but still reads the value versioned $(v_1 + k, p_1)$. In above case, the permutation of operations on p_0 and p_1 goes like this: (The version that a write assigns or a read returns are displayed in brackets.) $w(v_1 + 1, p_1) \prec_\pi w(v_1 + 2, p_1) \prec_\pi \dots \prec_\pi w(v_1 + k, p_1) \prec_\pi w(v_0 + 1, p_0) \prec_\pi w(v_0 + 2, p_0) \prec_\pi \dots \prec_\pi w(v_0 + k, p_0) \prec_\pi r(v_1 + k, p_1)$. Since $r(v_1 + k, p_1)$ returns a value not written by one of the latest k preceding writes in π , so the execution violates k -atomicity. By contradiction, we prove the correctness of Theorem 3.1. \square

3.2 Quantification of Visible Writes in One-Round-Write

W1R2 & W1R1 not only don't satisfy bounded k -atomicity, but also lead to frequent invisible writes that cannot be read. Here, we quantify the rate of visible writes incurred in W1R2 & W1R1 using the following assumption. Assume a Poisson process with parameter λ for the scenario of each client issuing a sequence of write/read operations, but *ignore* the duration of each operation. That is, each write/read is regarded as an instant event taking effect on the register at its starting moment. Then, we use the same client model as is applied in Section ?? . That is, clients are divided into writers and readers which means each client can initiate either writes or reads, but not both. Suppose there are n_w writers ($n_w > 1$), identified 0, 1, ... $n_w - 1$ separately, as well as an arbitrary number of readers. Then for writers, we have n_w independent, parallel write sequences, each satisfying a Poisson process with parameter λ . Assume that at the initial moment, all n_w writers are informed of the version (v_0, id) for some specific register that we focus on here. Let $N^i(t)$ be the number of write operations issued by writer i from the initial moment till time t , where $0 \leq i \leq n_w - 1$. According to the formula of Poisson probability, we have

$$\mathbb{P}\{N^i(t) = k\} = \frac{(\lambda t)^k e^{-\lambda t}}{k!}$$

Assume that writer i stores the local sequence number $k - 1$ at time t . Let $V^i(t, k)$ be a random variable denoting the visibility of the un-occurred write versioned $(v_0 + k, i)$ at

time t . Then, $V^i(t, k)$ has two status: invisible (denoted 0) and visible (denoted 1). $V^i(t, k) = 1$ occurs if and only if all other writers' local sequence number stayed lower than $(v_0 + k, i)$ at time t . Since write sequences are independent, we have

$$\begin{aligned} \mathbb{P}\{V^i(t, k) = 1\} &= \prod_{j=0}^{i-1} \mathbb{P}\{N^j(t) < k + 1\} \prod_{j=i+1}^{n_w-1} \mathbb{P}\{N^j(t) < k\}^{n_w-1-i} \\ &= \left(\sum_{j=0}^k \frac{(\lambda t)^j e^{-\lambda t}}{j!} \right)^i \left(\sum_{j=0}^{k-1} \frac{(\lambda t)^j e^{-\lambda t}}{j!} \right)^{n_w-1-i} \end{aligned}$$

Here, we quantify the rate when $k = 1$ at the expected arrival time $t = \frac{1}{\lambda}$, when the probability of $N^j(t) = 1$ is the highest. From the numerical analysis, in which we have chosen $\lambda = 10s^{-1}$, we observe that (see Fig. 6; the detailed results can be found in Table 3):...




Fig. 6. The probabilities of visible writes from writers with different identifiers ($\lambda = 10s^{-1}$, $t = 100ms$).




Fig. 7. The probabilities of visible writes from the writer with the minimum identifier with λ varying ($n_w = 10$).

3.3 Improvement of One-Round-Write Algorithms

Imagine changing the procedure of *update* function. When a replica receives an "update" message and finds the given version lower than it stores, then it would fill the acknowledgement with extra information, like the sequence number of the stored version and reply to the client. A client would have to wait for acknowledgement from a majority of replicas and

TABLE 3
Numerical results on the probabilities of **invisible** writes in one-round-write algorithms without extra information in acknowledgement($\lambda = 10s^{-1}, t = 100ms$).

Writer numbers	$id = 0$	$id = 1$	$id = 2$	$id = 3$	$id = 4$	$id = 5$	$id = 6$	$id = 7$	$id = 8$	$id = 9$
$n_w = 2$	0.632121	0.264241								
$n_w = 3$	0.864665	0.729329	0.458659							
$n_w = 4$	0.950213	0.900426	0.800852	0.601703						
$n_w = 5$	0.981684	0.963369	0.926737	0.853475	0.70695					
$n_w = 6$	0.993262	0.986524	0.973048	0.946096	0.892193	0.784386				
$n_w = 7$	0.997521	0.995042	0.990085	0.98017	0.96034	0.92068	0.84136			
$n_w = 8$	0.999088	0.998176	0.996352	0.992705	0.98541	0.97082	0.94164	0.883279		
$n_w = 9$	0.999665	0.999329	0.998658	0.997316	0.994633	0.989265	0.97853	0.957061	0.914122	
$n_w = 10$	0.999877	0.999753	0.999506	0.999013	0.998025	0.996051	0.992102	0.984204	0.968407	0.936814

TABLE 4
Numerical results on the probabilities(lower bound) of **invisible** writes in one-round-write algorithms if sequence number is included in acknowledgement($\lambda = 10s^{-1}, t = 100ms$).

Writer numbers	$id = 0$	$id = 1$	$id = 2$	$id = 3$	$id = 4$	$id = 5$	$id = 6$	$id = 7$	$id = 8$	$id = 9$
$n_w = 2$	0.264241	0.0803014								
$n_w = 3$	0.458659	0.323324	0.154154							
$n_w = 4$	0.601703	0.502129	0.377662	0.222077						
$n_w = 5$	0.70695	0.633687	0.542109	0.427636	0.284545					
$n_w = 6$	0.784386	0.730482	0.663103	0.578878	0.473598	0.341997				
$n_w = 7$	0.84136	0.8017	0.752125	0.690156	0.612695	0.515869	0.394836			
$n_w = 8$	0.883279	0.854099	0.817624	0.77203	0.715037	0.643796	0.554745	0.443431		
$n_w = 9$	0.914122	0.892652	0.865815	0.832269	0.790336	0.73792	0.6724	0.5905	0.488125	
$n_w = 10$	0.936814	0.921018	0.901272	0.87659	0.845738	0.807172	0.758965	0.698707	0.623383	0.529229

update its local sequence number according to the information in acknowledgement. Then, an invisible one-round-write may work like the *query* function and help update the local sequence number on the client side. In this way, W1R2 can be proved to satisfy n_w -atomicity given n_w writers.

Here we quantify the rate of visible writes in the improved W1R2 & W1R1 algorithms. Suppose writer i issues a write w after time t from its last write operation w' . Let $U^i(t)$ be a random variable denoting the visibility of the write w . Similarly, $U^i(t)$ has two status: invisible(denoted 0) and visible(denoted 1). Here, we only focus on one common pattern that would inevitably make w invisible: during the period from the finish time of w' to the invocation of w , there exists a writer identified smaller than i completing at least three writes, or a writer identified larger than i completing at least two writes. It can be proved through a case-by-case analysis. Since write sequences are independent, through probabilistic and combinatorial analysis we have

$$\begin{aligned}
& \mathbb{P}\{U^i(t) = 0\} \\
& \geq 1 - \prod_{j=0}^{i-1} \mathbb{P}\{N(t) < 3\} \prod_{j=i+1}^{n_w-1} \mathbb{P}\{N(t) < 2\} \\
& = 1 - e^{-\lambda t(n-1)} \cdot (1 + \lambda t + \frac{1}{2}(\lambda t)^2)^i \cdot (1 + \lambda t)^{n_w-1-i}
\end{aligned}$$

Then, we have

$$\begin{aligned}
& \mathbb{P}\{U^i(t) = 1\} = 1 - \mathbb{P}\{U^i(t) = 0\} \\
& \leq e^{-\lambda t(n-1)} \cdot (1 + \lambda t + \frac{1}{2}(\lambda t)^2)^i \cdot (1 + \lambda t)^{n_w-1-i}
\end{aligned}$$

Here, we quantify the rate at the expected time $t = \frac{1}{\lambda}$, which has the highest probability of that there exists no other operations between w' and w . From the numerical analysis, in which we have chosen $\lambda = 10s^{-1}$, we observe that (see Fig. 8; the detailed results can be found in Table 4):...




Fig. 8. The probabilities of visible writes from writers with different identifiers($\lambda = 10s^{-1}, t = 100ms$).

3.4 Brief Summary of One-Round-Write Algorithms

[1. 不保证有界 2. 无效写概率高]

From the counterexample of Section 3.1 we can see that, W1R2 and W1R1 don't promise to return value with bounded staleness only for *multiple-writer* registers⁵. Besides, the staleness of value returned for each read is related to operation distribution from different clients. What's more, the numeral results have shown that both algorithms don't satisfy atomicity most of the time, thus they cannot be available usually.

5. However, for *single-writer* multiple-reader registers, Wei [2] has proved that both algorithms guarantee providing deterministically **bounded** staleness of data versions for each read.




Fig. 9. The probabilities of visible writes from the writer with the minimum identifier by varying λ ($n_w = 10$).

4 Detailed Discussion of Experimental Results

4.1 ASC Tradeoff in Different Environments

We investigate whether the W2R1 algorithm can achieve the ASC tradeoff in different environments, by tuning workload pattern (including client number and read proportion), replica factor, communication delay and read quorum level.

4.1.1 Impact of Client Number

The client number reflects the concurrency patterns of operations. Intuitively, more clients will result in higher concurrency. In the experiment, each client is acted by one YCSB instance. All clients are allowed to initiate read/write requests concurrently, while each of them can only execute at most one request at any time. By varying client number from 10 to 40, we derive the results shown in Fig.10.

First we have a look at the consistency performance of each algorithm. W2R2 guarantees atomicity all the time as is expected, while other algorithms can lead to atomicity violations more or less. From the perspective of worst case we produced, W2R1 produced some non-atomic traces with k up to 4, while others results in a better performance. From the perspective of average consistency, W2R1 also has the worst performance compared with other three algorithms that applies Cassandra’s mechanisms (digest request, read repair or snitch strategy). However, although all but W2R2 can’t satisfy atomicity all the time, they guarantee atomicity most of the time, with a confidence that more than 99.97% read requests can obtain the most up-to-date value. In overall, the average consistency performance become weaker as the concurrent client number increases, and we produced the worst case of $k = 4$ when client number is up to 40. We speculate that the rate of returning fresh data for reads would decrease with more concurrent clients operating on single register, while all above non-atomic algorithms can also promise atomicity most of the time when concurrent clients are no more than 50, which is applicable for many applications and scenarios.

As for latency, we observe that average request latency is positively relative to client number in a subtle degree. The main reason is obvious: the processing capacity of each server is limited and would be gradually exhausted as the user throughput increases. From another aspect, W2R2 cost higher *read* latency compared with other four algorithms, while it gains stronger consistency all the time. Different algorithms

behaves in a slightly different way and we’ll discuss it in detail in Section 4.2.




Fig. 10. Consistency and latency results by tuning the number of clients.

4.1.2 Impact of Read Proportion

In the experiments, each client can initiate read and write requests. The ratio between read and write can also be a critical factor to impact the chance to produce violation patterns. Here, we use *read proportion* (the ratio of read numbers to total operation numbers) to describe the relations between read and write numbers. By varying read proportion from 0.50 to 0.99, we derive the results shown in Fig.11. As we can see, W2R2 still guarantees atomicity while others do not. W2R1 with snitch leads to a monotonically decreasing atomicity with higher read proportion. For W2R1, digest and repair algorithms, higher read proportion within certain range (0.5 - 0.9) can result in weaker consistency in average. However, when the read proportion is up to 0.99, the consistency performance reversely becomes better. One potential reason is that, k -atomicity violations are described by the staleness of values returned by reads, so within certain range, more reads will have more chances to form more complicated read-write concurrent patterns and return stale results. However, as the ratio of reads grows too high, the write will become too rare and sparsely distributed for reads to obtain stale written values.

As for latency, we observe that average request latency is approximately irrelative to read proportion for all algorithms except W2R1 with digests. The reason why request latency for W2R1 with digests grows higher when read proportion gets lower is mainly that, when write proportion becomes relatively higher, more frequent updates will make digest mismatch occur more frequently, triggering extra round-trip for collecting full data. If we only focus on *read* latency, W2R2 cost the highest *read* latency among all, while it guarantees strong consistency. Different algorithms behaves in a slightly different way and we’ll discuss it in detail in Section 4.2.




Fig. 11. Consistency and latency results by tuning read proportion.

4.1.3 Impact of Replica Factor

The replica factor in Cassandra specifies the number of replicas in each data center. The number of replicas impacts the consistency and latency performance in some degree, while the data center location of replicas impacts the request latency. Given a number of replicas, besides spreading them into different data centers, we can also place them all into a single data center for the convenience of local access. In the experiments, the replica factor is varied in 3 (inside single data-center), 1_1_1, 3_1_1, and 3_3_3. As is shown in Fig.12, more replicas would result in stronger consistency and higher latency in overall. The reason is obvious: the latency of an operation depends on the slowest responding replica which costs the longest time to communicate. In the network where delays are randomly distributed, it's more probable to have longer delays when communicating with more replicas. However, W2R1 *with snitch* behaves in a unique way. Applying smart routing through the snitch strategy makes the algorithm always process each read/write request by accessing local replicas in priority, which tends to lower the overall latency, but replica updates in remote data centers are usually not able to spread in time. 3_1_1 is a powerful evidence for above explanation.




Fig. 12. Consistency and latency results by tuning replica factor.

4.1.4 Impact of Network Delay

As is described before, to simulate a distributed storage system across WAN, we inject extra random delays in Cassandra to distinct replicas in different locations. Here, we mainly vary the inter-data center delay because it impacts the performance in a significant way. The injected inter-data center delays in network communication are normally distributed in the form of $N(\mu, \sigma)$, with $\mu = 100ms$ and σ varying from $5ms$ to $25ms$. From Fig. 13 we observe that, the average latency of write/read operations are basically in proportion to network delays among servers. Besides, the average consistency performance become weaker with the network jitters grow higher and higher. This is due to, not only higher network delay would make each operation duration stay longer, but also larger jitters or variances of delays would aggravate the occurrence of out-of-sync replicas. Then, the joint effects of above both lead to frequent atomicity violation patterns.

4.1.5 Impact of Read Quorum Level

We test three types of read quorum level in Cassandra: QUORUM, LOCAL_QUORUM, EACH_QUORUM. The level of QUORUM requires to access a quorum of replicas from all data-centers, which is actually the same as what we mention in theoretical model. The level of LOCAL_QUORUM only requires to access a quorum of replicas in the same




Fig. 13. Consistency and latency results by tuning inter-DC delay jitter.




Fig. 14. Consistency and latency results by tuning inter-DC delay.

data-center, which does actually not satisfy the quorum-replicated properties described in Lemma ???. However, without communicating with replicas in different data-centers, it speeds up requests a lot in sacrifice of consistency. The level of EACH_QUORUM requires to access a quorum of replicas from each data-centers, which is actually the stronger than the requirement of theoretical model. As is shown in Fig. ??, EACH_QUORUM promises the strongest average consistency performance among these three, and LOCAL_QUORUM the weakest. However, LOCAL_QUORUM has the lowest average latency, while EACH_QUORUM the highest. The tradeoff between consistency and latency is well presented in this experiment.




Fig. 15. Consistency and latency results by tuning read quorum level.

4.2 ASC Tradeoff with Different Algorithms

From all above experimental results by varying different factors, we observe that distinct algorithms vary in performance. Here, we summarize some insights from the following several perspectives.

- **Consistency** No matter Under what circumstance, W2R2 guarantees atomicity all the time, while other algorithms do not. Although from the staleness of returned data, other algorithms produced non-atomic traces with k up to 4 at worst, the probability of that is super low. From the frequency of atomicity violations, all algorithms

except W2R2 produce stale reads with the probability lower than 0.02% in our experiments under all circumstances. For example, from the experimental results in Table ..., the rate of 2-atomic reads is lower than 0.03%, and 3-atomic reads even lower, not to mention higher values of k in terms of k -atomic reads. Specifically, W2R1 with repair performs best except W2R2, and its probability of atomic reads in average is merely a little inferior than W2R2. W2R1 with digests or snitch don't perform as well as W2R1 with repair, but better than W2R1 in overall. Thus, we believe that all these algorithms rarely violate atomicity, and still promise atomicity most of the time.

- **Latency** We mainly focus on read latency. Normally, W2R1 presents lower read latency than W2R2, which is expected. Surprisingly, W2R1 with digests costs the largest read latency in average. We speculate in two reasons. Firstly, using digest requests means that each time when processing read requests, only one replica would be read with full data, and all others with digests, which cannot promise to read in only one round-trip once there exists some inconsistent replicas. Secondly, compared to raw data collection, digest requests reduces data size in communication but requires extra processing work, which takes extra time. W2R1 with repair leads to approximately the same or just slightly higher read latency compared to W2R1 without repair. This is due to that it applies asynchronous mechanism to repair stale replicas in background after the response of read operations, which has nearly no effects on the read response time. As for W2R1 with snitch, it gains the highest efficiency for reads mainly because it always selects the replicas with high proximity and route requests to them, which significantly decreases latency.

Among all algorithms, W2R2 gains the strongest consistency at the cost of extra communication round-trip of each read, which inevitably leads to high read latency. By omitting the second round-trip of read, W2R1 without any optimization gains lower latency but the worst consistency performance in average, although it still guarantees atomicity most of the time.

Compared with the original W2R2 or W2R1 algorithms, the optimization in Cassandra can either help increase consistency through replica synchronization(*repair*), or speedup request processing by reducing communication data size(*digests*) or smart routing(*snitch*). Applying these extra replica-related mechanisms in Cassandra to W2R1 separately, we observe that W2R1 with snitch has the lowest read latency due to its smart routing, but produces weakest consistency among all. W2R1 with repair performs better in promising atomicity without increasing much latency in average since it puts repair work in background. W2R1 with digests performs worse than W2R1 with repair but better than W2R1 with snitch in consistency, while its read latency is a lot higher than other one-round-trip read algorithms.

Among all, we think W2R1 with *repair* performs best concerning consistency and latency performance for two reasons. Firstly, the wisdom of this algorithm lies in the *divergence-oriented communication*. That is, only when inconsistency has been found will the algorithm take extra communication round-trip for repair, which costs less time for most consistent

requests while still keeps a high degree of atomicity. Secondly, repairing replicas in background asynchronously is another core technique for optimizing both consistency and latency performance.

For MWMR registers, one-round-write algorithms using Lynch's discrete version are far away from use. The results of W1R2 and W1R1 are presented in Fig. ?. As we can see, that most of the time both one-round-write algorithms cannot return an up-to-date value for a read, although they lowers write latency compared to that of two-round-write algorithms. However, Cassandra and many other storage systems use timestamps rather than Lynch's discrete version for data and apply one-round-write algorithms in production. In this way, no matter how precise the timestamp is, they cannot promise to always avoid producing identical timestamps for different data. Thus, it's out of scope in this research.

4.3 Results of One-Round-Write Algorithms

补：实验结果图

References

- [1] P. B. Gibbons and E. Korach, "Testing shared memories," *SIAM Journal on Computing*, vol. 26, no. 4, pp. 1208–1244, 1997.
- [2] H. Wei, Y. Huang, and J. Lu, "Probabilistically-atomic 2-atomicity: Enabling almost strong consistency in distributed storage systems," *IEEE Transactions on Computers*, vol. 66, no. 3, pp. 502–514, 2017.