Appendix to "Enabling Almost Strong Consistency for Quorum-replicated Datastores"

Lingzhi Ouyang, Yu Huang, Hengfeng Wei, Jian Lu

---

1 MONOTONICITY PROPERTIES OF SHARED REGISTER EMULATION ALGORITHMS

In this section, we study the properties of quorum-based algorithms theoretically. For simplicity, we assume that all operations in the history are applied to the same register. Some conventions are specified as follows:

- $\mathbb{W}$: the set of write operations.
- $\mathbb{R}$: the set of read operations.
- $\mathbb{O} = \mathbb{W} \cup \mathbb{R}$: the set of operations including writes and reads.
- $[o_{st}, o_{ft}]$: the time interval of the operation o. The time of invocation event and response event of o are denoted o_{st} and o_{ft} separately.
- $ver(o)$: the version attached to the returned value if operation $o \in \mathbb{R}$, or the version attached to the written value if $o \in \mathbb{W}$.
- $ver(s, t)$: the version stored in the server s at time t.
- $maxVer(A,t)/minVer(A,t)$: the maximum / minimum version stored in the server set A at a specific time t.
- $maxVer(A)/minVer(A)$: the maximum / minimum version collected in a query, where A is the responding server set of the query. Note that each server in A may respond not at the same time.

The pseudo-code for client-server interaction appears in Algorithm 1. It can be easily proved that:

Lemma 1.1 (Query-after-Update Property). *For any update u and any query q, if $u \prec_\sigma q$, and q receives responses from the server set Q, then $maxVer(Q) \geq ver(u)$.*

Proof. According to the server procedures in Algorithm 1, the version stored in servers keep non-decreasing as time goes by. Therefore, for any update u, when it is finished at time t_1, there exist a set of servers, denoted U, storing replicas with the version no lower than $ver(u)$. Namely, $minVer(U, t_1) \geq ver(u)$.

Now we consider any query q that starts after u. Namely, $u \prec_\sigma q$. Let Q denote the server set that q receives responses from.

According to the client procedures in Algorithm 1, any query or update will only be finished until a *majority* of servers reply, so $U \cap Q \neq \emptyset$. Assume the server $s \in U \cap Q$, and s returns the version at t_2 in the query q. Since q starts after u, then $t_2 > t_1$. Therefore, $maxVer(Q) \geq ver(s, t_2) \geq ver(s, t_1) \geq minVer(U, t_1) \geq ver(u)$. $\qquad\square$

The pseudo-codes for read/write operations appears in Algorithm 2 & 3. Here, we mainly focus on the properties in terms of monotonicity quorum-based algorithms, as shown in Table 1. These properties depict the relations between the temporal real-time relations and the semantic read-from relations of operations. The detailed proofs are mainly based on the *Query-after-Update Property*, and we display them in Section 1.1 - 1.4 respectively.

TABLE 1
Algorithm properties in terms of monotonicity

Properties	W2R2	W2R1	W1R2	W1R1
$w \prec_\sigma r \Rightarrow ver(w) \leq ver(r)$	✓	✓	✓	✓
$w \prec_\sigma w' \Rightarrow ver(w) < ver(w')$	✓	✓	×	×
$r \prec_\sigma r' \Rightarrow ver(r) \leq ver(r')$	✓	×	✓	×
$r \prec_\sigma w \Rightarrow ver(r) < ver(w)$	✓	×	×	×

Algorithm 1: Client-server interaction

```
1  ▷ Code for client process pᵢ(0 ≤ i ≤ n − 1):

2  function query (key)
3      vals ← ∅
4      pfor each server sⱼ              ▷ pfor: parallel for
5          send ['query', key] to sⱼ
6          v ←[key, val, ver] from sⱼ
7          vals ← vals ∪ v
8      until a majority of them respond
9      return vals

10 function update (key, value, version)
11     vals ← ∅
12     pfor each server sⱼ
13         send ['update', key, value, version] to sⱼ
14     wait for ['ACK']s from a majority of them

15 ▷ Code for server process sᵢ(0 ≤ i ≤ N − 1):

16 upon receive ['query', key] from pⱼ
17     send ['query − back', key, val, ver] to pⱼ

18 upon receive ['update', key, value, version] from pⱼ
19     pick [k, val, ver] with k == key
20     if ver < version then
21         val ← value
22         ver ← version
23     send ['ACK'] to pⱼ
```

Algorithm 2: Write algorithms for client p_i

---

1 procedure TwoRoundWRITE $(key, value)$
2     $replicas \leftarrow$ query(key)
3     $version \leftarrow$(maxSeq$(replicas)$ +1, $i)$
4   update$(key, value, version)$

5 procedure OneRoundWRITE $(key, value)$
6     $localSeq[key] \leftarrow localSeq[key]$ +1
7     $version \leftarrow(localSeq[key], i)$
8   update$(key, value, version)$

Algorithm 3: Read algorithms for client p_i

---

1 procedure TwoRoundREAD (key)
2     $replicas \leftarrow$ query(key)
3     $version \leftarrow$maxVer$(replicas)$
4     $value \leftarrow$valWithMaxVer$(replicas, version)$
5     $localSeq[key] \leftarrow version.seq$ $\triangleright$ Optional.
6   update$(key, value, version)$
7   **return** $value$

8 procedure OneRoundREAD (key)
9     $replicas \leftarrow$ query(key)
10     $version \leftarrow$maxVer$(replicas)$
11     $value \leftarrow$valWithMaxVer$(replicas, version)$
12     $localSeq[key] \leftarrow version.seq$ $\triangleright$ Optional.
13   **return** $value$

1.1 Common Property: Write-Read Monotonicity

The quorum-based write / read algorithms with one or two communication round-trips guarantee the *write-read monotonicity* for multi-writer registers. Specifically,

Theorem 1.1 (Write-Read Monotonicity). *In quorum-based algorithms(see Algorithm 1, 2, 3), for any operations w, r on the same register in the execution history σ, where $w \in \mathbb{W}, r \in \mathbb{R}$, if $w \prec_\sigma r$, then $ver(w) \leq ver(r)$.*

Proof. $\forall w \in \mathbb{W}, r \in \mathbb{R}$, let u denote the update process of w, q denote the query process of r, and Q denote the server set that q receives responses from. Then $ver(w) = ver(u), ver(r) = maxVer(Q)$.

If $w \prec_\sigma r$, then, it's trivial to have $u \prec_\sigma q$. By Lemma 1.1, $maxVer(Q) \geq ver(u)$. Therefore, $ver(w) = ver(u) \leq maxVer(Q) = ver(r)$. $\square$

1.2 Two-Round-Trip Write Property: Write-Write Monotonicity

The two-round-trip write algorithm guarantees the *write-write monotonicity* for multi-writer registers. Specifically,

Theorem 1.2 (Write-Write Monotonicity). *In quorum-based algorithms that completes each write operation in 2 communication round-trips (see TwoRoundWRITE in Algorithm 2), for any write operations w_1, w_2 on the same register in the execution history σ, if $w_1 \prec_\sigma w_2$, then $ver(w_1) < ver(w_2)$.*

Proof. For two-round-trip write operations w_1, w_2, let u_1 denote the update process of w_1, q_2 denote the query process of w_2, and Q_2 denote the server set that q_2 receives responses

from. Then $ver(w_1) = ver(u_1), ver(w_2) > maxVer(Q_2)$ (Note: w_2 will construct a larger version according to the largest version in Q_2).

If $w_1 \prec_\sigma w_2$, then, it's trivial to have $u_1 \prec_\sigma q_2$. By Lemma 1.1, $maxVer(Q_2) \geq ver(u_1)$. Therefore, $ver(w_1) = ver(u_1) \leq maxVer(Q) < ver(w_2)$. $\square$

In comparison, the one-round-write algorithm can only guarantee the *write-write monotonicity* for *single-writer* registers, where the only writer can execute updates with *monotonically increasing local versions*. However, for *multi-writer* registers, one-round-write may result in *write-version inversion anomalies*, which means that a *previous* write assigns a *larger* version for its written value than a *later* write (Namely, $(w_1 \prec_\sigma w_2) \cap (ver(w_1) > ver(w_2))$). What's worse, the procedure without the promise of assigning a version larger than previous writes may lead to unsuccessful updates on servers. More details will be discussed in Section 3 of this appendix.

1.3 Two-Round-Trip Read Property: Read-Read Monotonicity

The two-round-trip read algorithm guarantees the *read-read monotonicity* for multi-writer registers. Specifically,

Theorem 1.3 (Read-Read Monotonicity). *In quorum-based algorithms that completes each read operation in 2 communication round-trips (see TwoRoundREAD in Algorithm 3), for any read operations r_1, r_2 on the same register in the execution history σ, if $r_1 \prec_\sigma r_2$, then $ver(r_1) \leq ver(r_2)$.*

Proof. For two-round-trip read operations r_1, r_2, let u_1 denote the update process of r_1, q_2 denote the query process of r_2, and Q_2 denote the server set that q_2 receives responses from. Then $ver(r_1) = ver(u_1), ver(r_2) = maxVer(Q_2)$.

If $r_1 \prec_\sigma r_2$, then, it's trivial to have $u_1 \prec_\sigma q_2$. By Lemma 1.1, $maxVer(Q_2) \geq ver(u_1)$. Therefore, $ver(r_1) = ver(u_1) \leq maxVer(Q) = ver(r_2)$. $\square$

In comparison, the one-round-read algorithm can guarantee the *read-read monotonicity* for single-reader registers only if clients store previous read records locally. However, for multi-reader registers, the one-round-read algorithm may lead to *version inversion anomalies after a read*, which means that the version of a later operation (a read or a write) is smaller than the version of a *previous* read. The anomalies involves several patterns and more details will be discussed in Section 2 of this appendix.

1.4 Exclusive Property of W2R2: Read-Write Monotonicity

Besides all above properties, the W2R2 algorithm also guarantees the *read-write monotonicity* for MWMR registers.

Theorem 1.4 (Read-Write Monotonicity). *In the W2R2 algorithm, for any operations w, r on the same register in the execution history σ, where $w \in \mathbb{W}, r \in \mathbb{R}$, if $r \prec_\sigma w$, then $ver(r) < ver(w)$.*

Proof. $\forall w \in \mathbb{W}, r \in \mathbb{R}$, let u_r denote the update process of r, q_w denote the query process of w, and Q_w denote the server set that q receives responses from. Then $ver(r) =$

$ver(u_r), ver(w) > maxVer(Q_w)$ (Note: w will construct a larger version according to the largest version in Q_w).

If $w \prec_\sigma r$, then, it's trivial to have $u_w \prec_\sigma q_w$. By Lemma 1.1, $maxVer(Q_w) \geq ver(u_r)$. Therefore, $ver(r) = ver(u_r) \leq maxVer(Q_w) < ver(w)$. $\qquad\square$

2 ATOMICITY VIOLATION OF W2R1: STALENESS AND PROBABILITY

In this section, we first prove the possible atomicity violation patterns in W2R1 (see Algorithm 4), then we prove the bound of data staleness and calculate the probability of atomicity violation in W2R1 based on the violation patterns.

Algorithm 4: W2R1

---

1 Code for client process $p_i(0 \leq i \leq n - 1)$:

2 **procedure** TwoRoundWRITE ($key, value$)

3 $replicas \leftarrow$ query (key)

4 $version \leftarrow$(maxSeq ($replicas$) +1,i)

5   update ($key, value, version$)

6 **procedure** OneRoundREAD (key)

7 $replicas \leftarrow$ query (key)

8 $version \leftarrow$maxVer ($replicas$)

9 $value \leftarrow$valWithMaxVer ($replicas, version$)

10   **return** $value$

2.1 Proof of Atomicity Violation Patterns

We aim to prove that atomicity violation incurred in W2R1 are composed of RI or WI. When clients read and write a MWMR register using the W2R1 algorithm and obtain the history σ, we have that:

Theorem 2.1. *If σ violates atomicity, then there exists some operations in σ that form either RI or WI.*

Proof. If σ violates atomicity, then for any permutation π of σ, we have a stale read r, the dictating write w of r and the interfering write w' satisfying: $w \prec_\pi w' \prec_\pi r$. Then, we exhaustively check all cases that make r a stale read.

According to the definition of atomicity, two types of relations between operations are of our concern: the temporal real-time relation and the semantic read-from relation. We first enumerate all possible cases according to the semantic relation. Then for each case, we further enumerate all possible sub-cases according to the temporal relation.

Since w is the dictating write of r, we have that $ver(r) = ver(w)$. According to the semantic relation between versions of w and w', we have two complementing cases: $ver(r) = ver(w) < ver(w')$ and $ver(r) = ver(w) > ver(w')$. In each case, we then consider the temporal relation between operations.

Case 1: $ver(r) = ver(w) < ver(w')$.

Note that r *must be concurrent with w'*. Namely, $r \parallel_\sigma w'$. This can be proved by contradiction. If $r \prec_\sigma w'$, w' can never be the interfering write of r. If $w' \prec_\sigma r$, according to the *write-read monotonicity* (Theorem 1.1), r must return the version of w' (or return a even larger version), contradicting the fact that $ver(r) = ver(w) < ver(w')$.

Next we exhaustively check all possible sub-cases according to the temporal relation to prove that: *there must exist a read operation r' that reads from w', and r' precedes r.* Namely, $\exists r' = R(w') : r' \prec_\sigma r$.

We first enumerate all possible cases according to the temporal relation between w and w'. Since in the permutation π, $w \prec_\pi w'$, we have that in σ, $w \prec_\sigma w'$ or $w\parallel_\sigma w'$.

Case 1.1: $w \prec_\sigma w'$. The permutation between w and w' must be *determined* as $w \prec_\pi w'$. Then, we enumerate all cases as follows:

$\times$ $\nexists r' = R(w')$. Note that $w' \parallel_\sigma r$, so r can be ordered before w'. Thus, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, which contradicts the permutation $w \prec_\pi w' \prec_\pi r$.

$\times$ $\forall r' = R(w') : r' \nprec_\sigma r$. Similarly, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, contradicting the permutation $w \prec_\pi w' \prec_\pi r$.

$\checkmark$ $\exists r' = R(w') : r' \prec_\sigma r$. Then the permutation among w, w' and r must be $w \prec_\pi w'(\prec_\pi r') \prec_\pi r$ (see Fig. 1 (a)).

Case 1.2: $w \parallel_\sigma w'$. In this case, the permutation between w and w' can not be *determined* only by the temporal relation of these two write operations.

Therefore, we then focus on other operations related to these two clusters [1]. We first consider the temporal relation between w' and any read operation r'' that reads from w. Notice that r is a special instance of r'', thus r'' must exist.

Similarly, by contradiction we have $w' \nprec_\sigma r''$; otherwise r'' must return the version of w' (or return a even larger version) according to the *write-read monotonicity* (Theorem 1.1). Thus, the temporal relation between w' and r'' can only be $r'' \prec_\sigma w'$ or $r'' \parallel_\sigma w'$.

Case 1.2.1: $\exists r'' = R(w) : r'' \prec_\sigma w'$. By the temporal relation between r'' and w' , as well as read-from order between r'' and w, the permutation between w and w' must be *determined* as $w(\prec_\pi r'') \prec_\pi w'$.

$\times$ $\nexists r' = R(w')$. Note that $w' \parallel_\sigma r$, so r can be ordered before w'. Thus, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, contradicting the permutation $w \prec_\pi w' \prec_\pi r$.

$\times$ $\forall r' = R(w') : r' \nprec_\sigma r$. Similarly, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, contradicting the permutation $w \prec_\pi w' \prec_\pi r$.

$\checkmark$ $\exists r' = R(w') : r' \prec_\sigma r$. Then the permutation must be $w \prec_\pi w'(\prec_\pi r') \prec_\pi r$ (see Fig. 1 (b)).

Case 1.2.2: $\forall r'' = R(w) : r'' \parallel_\sigma w'$. Then, We consider the temporal relation between w and any read operation r' that reads from w'. Note that r' may not exist.

Case 1.2.2.1: $\nexists r' = R(w')$. Then, there exists no determining factor to deciding the permutation between w and w'. Thus, the permutation among w, w', r is $w' \prec_\pi w \prec_\pi r$

---

1. Here, we use the concept *cluster*, a terminology introduced by Gibbons and Korach [1], for our analysis. A *cluster* is a subset of operations in an execution that consists of a write and all of its dictated reads. Gibbons and Korach [1] have shown that atomicity violation indicates certain rules related to different zones of clusters. Therefore, for soundness of proof, atomicity violation concerned with w, w', r should be analyzed with all operations of their related clusters, consisting of w and all of its dictated reads, as well as w' and all of its dictated reads, because any dictating reads that returns related written values may effect the permutation.

or $w \prec_\pi r \prec_\pi w'$, both contradict the assumption that w, w' and r form an inconsistent read. [×]

Case 1.2.2.2: $\exists r' = R(w') : r' \prec_\sigma w$. By the temporal relation between r' and w , as well as read-from order between r' and w', the permutation between w and w' is *determined* as $w'(\prec_\pi r') \prec_\pi w$,which contradicts $w \prec_\pi w'$. [×]

Case 1.2.2.3: $\forall r' = R(w') : r' \not\prec_\sigma w$. Then, there exists no determining factor to deciding the permutation between w and w'.

× $\forall r' = R(w') : r' \not\prec_\sigma r$. Thus, the permutation among w, w' and r is $w' \prec_\pi w \prec_\pi r$ or $w \prec_\pi r \prec_\pi w'$, both violating the given permutation $w \prec_\pi w' \prec_\pi r$.

✓ $\exists r' = R(w') : r' \prec_\sigma r$. Then the permutation can be $w \prec_\pi w' \prec_\pi r' \prec_\pi r$ or $w' \prec_\pi w \prec_\pi r' \prec_\pi r$ (see Fig. 1 (c)), which involves atomicity violation anyway.

Above all, for $ver(w) < ver(w')$, then w, w' and r can be linearly extended to $w \prec_\pi w' \prec_\pi r$ only when $\exists r' = R(w') : r' \prec_\sigma r$. Obviously, r and r' form RI.


Fig. 1. Typical patterns of RI in Case 1.

Case 2: $ver(r) = ver(w) > ver(w')$.

Note that w must be concurrent with w'. This can be proved by contradiction. If $w \prec_\sigma w'$, according to *write-write monotonicity* that two round-trip write algorithm holds (Theorem 1.2), the version of w' must be larger. If $w' \prec_\sigma w$, we can never get $w \prec_\pi w' \prec_\pi r$ when linearly extending σ to π. Besides, according to the permutation $w' \prec_\pi r$, we have $w' \prec_\sigma r$ or $w' \parallel_\sigma r$.

Next we enumerate all possible sub-cases according to the temporal relation to prove that: there exists a read operation r'' (can be r) that reads from w, satisfying that r'' precedes either w' or another read operation r' that reads from w'. Namely, $\exists r'' = R(w) : (r'' \prec_\sigma w') \vee (\exists r' = R(w') : r'' \prec_\sigma r')$.

We first consider the temporal relation between w' and any read operation r'' that reads from w. Notice that r is a special instance of r'', thus r'' must exist.

Case 2.1: $\exists r'' = R(w) : r'' \prec_\sigma w'$. Then, the permutation between w and w' must be *determined* as $w(\prec_\pi r'') \prec_\pi w'$. Note that r cannot take the role of r'' here, because $r \not\prec_\sigma w'$.

Case 2.1.1: $w' \prec_\sigma r$. Then the permutation among w, w' and r must be $w \prec_\pi w' \prec_\pi r$ (see Fig. 2 (a)). [✓]

Case 2.1.2: $w' \parallel_\sigma r$.

× $\nexists r' = R(w')$. Note that $w' \parallel_\sigma r$, so r can be ordered before w'. Thus, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, contradicting the assumption.

× $\forall r' = R(w') : r' \not\prec_\sigma r$. Similarly, the permutation among w, w' and r is $w \prec_\pi r \prec_\pi w'$, contradicting the assumption.

✓ $\exists r' = R(w') : r' \prec_\sigma r$. Then the permutation among w, w' and r must be $w(\prec_\pi r'') \prec_\pi w'(\prec_\pi r') \prec_\pi r$ (see Fig. 2 (b)).

Case 2.2: $\forall r'' = R(w) : r'' \not\prec_\sigma w'$. Then, We consider the temporal relation between r'' and any read operation r' that reads from w'.

Case 2.2.1: $\nexists r' = R(w')$. Then, there exists no determining factor to deciding the permutation between w and w'.

× $w' \prec_\sigma r$. Then the permutation among w, w' and r must be $w' \prec_\pi w \prec_\pi r$, contradicting the assumption.

× $w' \parallel_\sigma r$. Thus, the permutation among w, w', r is $w' \prec_\pi w \prec_\pi r$ or $w \prec_\pi r \prec_\pi w'$, both contradicting the assumption.

Case 2.2.2: $\forall r' = R(w'), r'' = R(w) : r'' \not\prec_\sigma r'$.

× $w' \prec_\sigma r$. Then the permutation among w, w' and r must be $w' \prec_\pi w \prec_\pi r$, contradicting the assumption.

× $w' \parallel_\sigma r$. Thus, the permutation among w, w', r is $w' \prec_\pi w \prec_\pi r$ or $w \prec_\pi r \prec_\pi w'$, both contradicting the assumption.

Case 2.2.3: $\exists r' = R(w'), r'' = R(w) : r'' \prec_\sigma r'$.

By contradiction we have $w \not\prec_\sigma r'$; otherwise r' must return the version of w (or return a even larger version) according to the *write-read monotonicity* (Theorem 1.1). Besides, r'' can't precede its dictating write w due to the regularity property. Thus, the temporal relation can only be $w \parallel_\sigma r''$ as well as $w \parallel_\sigma r'$. The permutation among w, r'', r' must be $w \prec_\pi r'' \prec_\pi r'$. Then the permutation can be $w \prec_\pi w' \prec_\pi r'' \prec_\pi r'$ or $w' \prec_\pi w \prec_\pi r'' \prec_\pi r'$, which involves atomicity violation anyway. If $w' \prec_\sigma r$ (or r just takes the role of r''), then r may become a stale read (see Fig. 2 (c)). [✓]

Above all, for $ver(w) > ver(w')$, then w, w' and r can be linearly extended to $w \prec_\pi w' \prec_\pi r$ only when there exists a read operation r'' (can be r) that reads from w, satisfying that r'' precedes either w' or another read operation r' that reads from w'. Namely, $\exists r'' = R(w) : (r'' \prec_\sigma w') \vee (\exists r' = R(w') : r'' \prec_\sigma r')$, where r'' and w' form WI, or r'' and r' form RI.


Fig. 2. Typical patterns of WI (Sub-fig (a) and (b))and RI (Sub-fig (c)) in Case 2.

In conclusion, by assuming r is a stale read, we exhaustively check all cases according to the temporal real-time relations and the semantic read-from relations of relative operations, At last, we have proved that the necessary conditions to make r a stale read involve RI or WI. Therefore, Theorem 2.1 is true. □

From the proof of Theorem 2.1 we can obtain some more detailed information, which is useful for our further analysis of staleness bound.

Proposition 2.1. *Suppose σ is the history obtained when clients read and write a MWMR register using the W2R1 algorithm. For any permutation π of σ, there exists a stale read r, the dictating write w of r and the interfering write w' that form $w \prec_\pi w' \prec_\pi r$, satisfying **Case 1** or **Case 2**, where*

- **Case 1:** $ver(w) < ver(w')$. w' and r are concurrent, and there exists a read r' that reads from w', as well as precedes r. Namely, $(w' \parallel_\sigma r) \land (\exists r' = R(w') : r' \prec_\sigma r)$;
- **Case 2:** $ver(w) > ver(w')$. w and w' are concurrent, and there exists a read r'' (can be r) that reads from w, satisfying that r'' precedes either w' or another read r' that reads from w'. Namely, $(w \parallel_\sigma w') \land (\exists r'' = R(w) : (r'' \prec_\sigma w') \lor (\exists r' = R(w') : r'' \prec_\sigma r'))$.

2.2 Proof of Bound of Data Staleness

In this section, we calculate the tight bound of data staleness when accessing a MWMR register using the W2R1 algorithm. Suppose the distributed storage system consists of N replicas($N \geq 3$). Denote the number of writer clients as n_w. For any history σ, we have that:

Theorem 2.2. *There exists a linear extension π of σ such that any read in π returns the value of the latest B proceeding write. Here, $B = n_w + \frac{1}{2}n_w(n_w - 1) + 1$. Moreover, the bound B is tight, i.e. there exists a permutation π in which some read returns values of the oldest write in the latest B preceding writes.*

Proof. The proof of Theorem 2.2 is based on an adversary argument, to insert as many interfering writes as possible before the inconsistent read. The proof also needs to consider the same two cases as defined in Section 2.1. Specifically, suppose r is an inconsistent read, and the dictating write of r is w. There must exist another interfering write w' such that $w \prec_\pi w' \prec_\pi r$.

According to the versions $ver(r)$ and $ver(w')$, we consider two cases. In Case 1 where $ver(r) = ver(w) < ver(w')$, we prove that there are at most $B_1 = n_w$ interfering writes which can be inserted between r and w. In Case 2 where $ver(r) = ver(w) > ver(w')$, we prove that there are at most $B_2 = \frac{1}{2}n_w(n_w - 1)$ interfering writes.

Case 1: $ver(w) = ver(r) < ver(w')$.

From Case 1 of Proposition 2.1 we know that the interfering write w' must be concurrent with r, and there exists a read operation r' preceding r and dictated by w'. Note that the invocation of w' must be *earlier* than that of r, or r' is not able to read from w' when preceding r . Since each writer client can have at most one write operation that starts before r and is concurrent with r, the number of interfering writes in this case is no more than n_w. Thus we have the bound $B_1 = n_w$ in this case, and the bound B_1 is tight. One construction of most stale read is illustrated in Fig. 3, where the pattern of operations is shown in Fig. 1 (a).
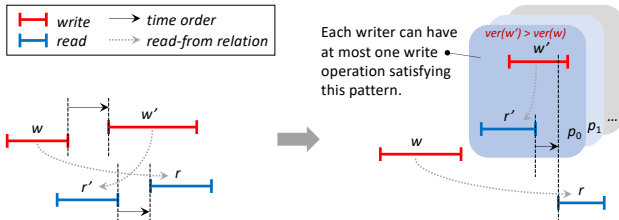

Fig. 3. Maximum number of interfering writes in Case 1.

Case 2: $ver(w) = ver(r) > ver(w')$.

From Case 2 of Proposition 2.1 we know that the interfering write w' must be concurrent with w. Other than the writer client of operation w, we can have at most $n_w - 1$ writer clients to generate interfering writes. The challenge in the construction is that, for each interfering writer client, we cannot insert as many interfering write operations w' as we want, due to the constraint of $ver(w') < ver(w)$ and $w' \parallel_\sigma w$.

Denote the write operation versioned (seq, i) as $w(seq, i)$. In this case, both WI and RI are able to coexist to result in stale reads, we need to consider both patterns in our adversary argument. We first construct a simplified worst-case trace using WI alone, then we take both patterns into consideration.

The construction of WI without RI

In WI, there exists a read r'' that reads from $w(seq, i)$ and precedes w'. We focus on the maximum number of interfering write operations w' that can be inserted after r'' while still concurrent with $w(seq, i)$. Let w have the version (seq, i). The key point of our adversary argument is to keep the *maximum* version of some replica majority as *small* as possible before $w(seq, i)$ is done, so the interfering writes can have the chance to query that replica majority and write with the version value smaller than (seq, i).

Now let's describe the construction with adversary argument in detail. First of all, at the finish time of r'', if the maximum version of any replica majority become no smaller than (seq, i), then no interfering two-round-trip write operation after r'' can have the version value smaller than (seq, i). Therefore, there still exists some replica majority whose maximum version value is smaller than (seq, i) after the finish time of r'', and the interfering two-round-trip write operations may still have the chance to have the version value smaller than (seq, i).

Be aware that w won't have $ver(w) = (seq, i)$ unless it queries a majority of replicas whose maximum version is $(seq - 1, i_1)$ in its first round-trip. After the finish time of r'', if the maximum version of any replica majority becomes no smaller than $(seq - 1, i_1)$, then any interfering two-round-trip write operation can only have the version no smaller than (seq, i'), where $i' < i$. However, when there still exists some replica majority whose maximum version is smaller than $(seq-1, i_1)$ (the write operation $w(seq-1, i_1)$ is not yet finished at this time for sure), take $(seq - 2, i_2)$ for instance, then the interfering two-round-trip write operations may still have the chance to have the version smaller than $(seq - 1, i'_1)$. Since $(seq-1, i'_1) < (seq, i')$, for adversary argument, let $w(seq - 1, i_1)$ be unfinished at the finish time of r'', so more interfering write operations can be inserted after r''.

Recursively, be aware that a write operation won't have the version value $(seq - 1, i_1)$ unless it queries a majority of replicas whose maximum version is $(seq - 2, i_2)$ in its first round-trip. For adversary argument, let $w(seq - 2, i_2)$ be unfinished at the finish time of r'', so more interfering write operations can be inserted after r''. Similarly, let $w(seq - 3, i_3), w(seq-4, i_4), ..., w(seq-n+1, i_{n-1})$ be unfinished at the finish time of r'' until the operations of all writer clients (except the writer client of operation $w(seq, i)$) are in.

Above all, we adversely make a construction that allows the most interfering write operations in the pattern of WI. Note that all of $w(seq - 1, i_1)$, $w(seq - 2, i_2)$, ..., $w(seq - n + 1, i_{n-1})$ can be ordered before r'' in π. Then, after $w(seq - n + 1, i_{n-1})$ is finished while $w(seq - 1, i_1)$,

$w(seq-2, i_2)$, ..., $w(seq-n+2, i_{n-2})$ are processing, the writer client i_{n-1} may query a majority of replicas whose maximum version value is $w(seq-n+1, i_{n-1})$ and can issue an interfering write operation $w(seq-n+2, i_{n-1})$. Similarly, after $w(seq-n+1, i_{n-1})$ and $w(seq-n+2, i_{n-2})$ are finished while $w(seq-1, i_1)$, $w(seq-2, i_2)$, ..., $w(seq-n+3, i_{n-3})$ are processing, the writer clients i_{n-1} and i_{n-2} may query a majority of replicas whose maximum version value is $w(seq-n+2, i_{n-2})$, and then issue interfering write operations $w(seq-n+3, i_{n-1})$ and $w(seq-n+3, i_{n-2})$ respectively. In this way, after all of $w(seq-1, i_1)$, $w(seq-2, i_2)$, ..., $w(seq-n+1, i_{n-1})$ are done, all writer clients except i may query a majority of replicas whose maximum version value is $w(seq-1, i_{n-1})$, and then issue interfering write operations $w(seq, i')$, $i' < i$. Let i be the maximum client identifier. Then, the maximum number of interfering write operations using WI is:

$$B_2^{WI} = 1 + 2 + \cdots + (n_w - 2) + (n_w - 1) = \frac{1}{2} n_w(n_w - 1)$$

Fig. 4 shows a concrete example of the most stale read with WI when $n_w = 4$, where the pattern of operations is shown in Fig. 2 (a). The read operation r returns the value written by $w(100, 3)$ issued by the writer client with the maximum identifier.


Fig. 4. Maximum number of interfering writes in the pattern of WI in Case 2.

The construction of WI and RI

Now we discuss the situation with the additional effects of RI. For RI, according to Case 2 of Proposition 2.1, there exists a read r'' (can be r) that's dictated by w preceding another read r' dictated by w'. For adversary argument, let r be the read operation with the latest starting time dictated by w, so it must be the most stale read among those dictated by w.

Consider any write operation w' that's concurrent with w and versioned smaller than $ver(w)$. As for the dictating read r'' of w, we mainly focus on the one that *ends earliest*, because the position of w in π is largely decided by it compared to other dictating reads of w. Note that if w' is an interfering write of r, then either w' forms WI with r'', or a dictated read of w' forms RI with r''. Without losing generality, w' can be preceding, concurrent with or starting after r''.

Now we describe an approach to linearly extending σ to a permutation π where the number of interfering writes in Case 2 is no more than B_2^{WI}. For any w' that's preceding or concurrent with r'', let w' precede w and r'' in π, preventing it from interfering the dictating reads of w. For any w' that starts after r'', w' must be ordered after w and r''. Thus, each of these w' will interfere the dictating read(s) of w that start after it. Repeat above procedures recursively for all write

operations. In this way, all w' that start not after r'' won't interfere the dictating reads of w, and only those w' that starts after r'' can have the chance to interfere the dictating reads of w.

We now analyze the maximum staleness of each read dictated by w' or w in above permutation. For any read r' dictated by w', if it starts after some read operation r''' dictated by w, then r' and r''' form RI. Since $ver(r') = ver(w') < ver(w)$, according to the worst-case construction using RI in Case 1, the interfering write operations of r' is no more than $B_1 = n_w$. For any read r''' dictated by w, if it starts after some write w', then w' and r''' form the pattern of WI. According to the worst-case construction using WI in Case 2, the number of write operations versioned smaller than $ver(w)$ and starting after r'' is no more than $B_2^{WI} = \frac{1}{2} n_w(n_w - 1)$. Thus, the dictating read r''' of w can be guaranteed to return a value no more stale than B_2^{WI}.

Note that this approach only promises to obtain a legal permutation π of any given σ. Whether there exists some other legal permutation with smaller maximum staleness of reads or not, the bound of staleness in Case 2 is definitely not larger than B_2^{WI}.

Since the trace with B_2^{WI} interfering writes in Case 2 can be constructed with WI, we have the bound $B_2 = B_2^{WI} = \frac{1}{2} n_w(n_w - 1)$, and the bound B_2 is tight.

Given a specific read operation r, the interfering write operations in Case 1 and Case 2 can appear at the same time in the history σ, so we can construct a trace with $B_1 + B_2$ interfering writes. Counting the dictating write itself, we have that the read always returns the value of the latest

$$B = B_1 + B_2 + 1 = n_w + \frac{1}{2} n_w(n_w - 1) + 1$$

preceding writes, i.e. the history σ always satisfies B-atomicity. During the proof, we explicitly construct the worst-case trace, which contains the read having $B_1 + B_2$ interfering writes. This proves that the bound is tight. $\square$

An example of the worst-case construction is illustrated in Fig. 5. The process of version updates on replicas is shown in Table 2.

Till now, we prove the correctness of Theorem 2.2. The bound of data staleness using the W2R1 algorithm is related to the number of writer clients. Specially, for *single-writer* multiple-reader registers, W2R1 satisfies 2-atomicity which conforms to the conclusion in [2].

2.3 Probability of Atomicity Violation

In this subsection we quantify the atomicity violation incurred in W2R1. The quantification follows from theorem 2.1 that, atomicity violation incurred in W2R1 are composed of either RI or WI. We simplify the quantification problem into the following one: what's the probability that RI or WI incurred in W2R1? The simplification helps us to quantify atomicity violation by ignoring unnecessary details. Note that the occurrence of either RI or WI is merely the condition *necessary* but *not sufficient* for atomicity violation in W2R1, as

Fig. 5. An example of worst case with 3 writer clients and 3 reader clients using the W2R1 algorithm. The *version* that a *write* assigns or a *read* returns are displayed in brackets. One permutation of these operations can be: $w_0(98,0) \prec_\pi w_1(99,1) \prec_\pi \mathbf{w_2(100,2)} \prec_\pi r_5''(100,2) \prec_\pi w_0'(99,0) \prec_\pi w_0'(100,0) \prec_\pi w_1'(100,1) \prec_\pi w_0'(101,0) \prec_\pi r_3'(101,0) \prec_\pi w_1'(101,1) \prec_\pi r_4'(101,1) \prec_\pi w_1'(101,2) \prec_\pi r_3'(101,2) \prec_\pi \mathbf{r_5(100,2)}$. For $r_5(100,2)$, it returns the value of the latest 7 preceding write in π. Interfering write operations w' of r include $w_0'(99,0), w_0'(100,0), w_1'(100,1), w_0'(101,0), w_1'(101,1), w_2'(101,2)$. Among these operations, $r_5''(100,2)$ and each of $w_0'(99,0), w_0'(100,0), w_1'(100,1)$ form the pattern of WI respectively; $r_5(100,2)$ and each of $r_3'(101,0), r_4'(101,1), r_3'(101,2)$ form the pattern of RI respectively.

TABLE 2
The process of version updates on replicas in Fig. 5. Suppose the distributed storage system contains 3 replicas.

Time	Replica 1	Replica 2	Replica 3	Description
t_0	**(97,0)**	**(97,0)**	(97,0)	p_0 then assigns $(98,0)$ after the query returns maximum version $(97,0)$ in 1^{st} round-trip.
t_1	**(98,0)**	(97,0)	(97,0)	p_1 then assigns $(99,1)$ after the query returns maximum version $(98,0)$ in 1^{st} round-trip.
t_2	**(99,1)**	(97,0)	(97,0)	p_2 then assigns $(100,2)$ after the query returns maximum version $(99,1)$ in 1^{st} round-trip.
t_3	**(100,2)**	(97,0)	(97,0)	p_5 returns the value versioned $(100,2)$ in the read.
t_4	(100,2)	**(98,0)**	**(97,0)**	p_0 then assigns $(99,0)$ after the query returns maximum version $(98,0)$ in 1^{st} round-trip.
t_5	(100,2)	**(99,0)**	(97,0)	p_0 then assigns $(100,0)$ after the query returns maximum version $(99,0)$ in 1^{st} round-trip.
t_6	(100,2)	**(99,0)**	(97,0)	p_1 then assigns $(100,1)$ after the query returns maximum version $(99,0)$ in 1^{st} round-trip.
t_7	(100,2)	**(100,2)**	(97,0)	p_0 assigns $(101,0)$ after the query returns maximum version $(100,2)$ in 1^{st} round-trip.
t_8	(100,2)	**(101,0)**	(97,0)	p_3 returns the value versioned $(101,0)$ in the read.
t_9	**(100,2)**	(101,0)	(97,0)	p_1 assigns $(101,1)$ after the query returns maximum version $(100,2)$ in 1^{st} round-trip.
t_{10}	(100,2)	**(101,1)**	(97,0)	p_4 returns the value versioned $(101,1)$ in the read.
t_{11}	**(100,2)**	(101,1)	(97,0)	p_2 assigns $(101,2)$ after the query returns maximum version $(100,2)$ in 1^{st} round-trip.
t_{12}	(100,2)	**(101,2)**	(97,0)	p_3 returns the value versioned $(101,2)$ in the read.
t_{13}	**(100,2)**	(101,2)	**(97,0)**	p_5 returns the value versioned $(100,2)$ in the read.

shown in Fig. 6), but this will not prevent us from obtaining the upper bound of the violation probability:

$$
\begin{aligned}
\mathbb{P}\{\text{Atomicity Violation in W2R1}\} &\leq \mathbb{P}\{\text{RI} \vee \text{WI}\} \\
&= \mathbb{P}\{\text{RI}\} + \mathbb{P}\{\text{WI}\} - \mathbb{P}\{\text{RI} \wedge \text{WI}\} \\
&\leq \mathbb{P}\{\text{RI}\} + \mathbb{P}\{\text{WI}\}
\end{aligned}
\tag{2.1}
$$

Fig. 6. The relation between atomicity violation and two types of version inversions in W2R1 .

According to the case-by-case proof of theorem 2.1, RI and WI can only occur within certain conditions in W2R1. Intuitively, an execution with higher degree of concurrency would produce more atomicity violation, but it's still not sufficient without certain read-write patterns. Thus, we decompose RI and WI incurred in W2R1 into a *concurrency pattern* and a *read-write pattern* separately. The *concurrency pattern(CP)* describes the constraints of real-time orders among events of operations, and the *read-write pattern(RWP)* limits read/write semantics among operations. Then, we have

Definition 2.1. *The read version inversion after a read in W2R1 (RI-W2R1) are composed of one write w' and two reads r, r', satisfying the requirements of*

- *the long-lived-write concurrency pattern(CP):*
  1) $r_{st} \in [w'_{st}, w'_{ft}]$,
  2) $r'_{ft} \in [w'_{st}, r_{st}]$;
- *the non-monotonic read-write pattern(RWP):*
  3) $r' = R(w')$;
  4) $r \neq R(w')$.

Definition 2.2. *The write-version inversion after a read in W2R1 (WI-W2R1) are composed of two writes w, w' and one read r'', satisfying the requirements of*

- *the long-lived-write concurrency pattern(CP):*
  1) $w'_{st} \in [w_{st}, w_{ft}]$,
  2) $r''_{ft} \in [w_{st}, w'_{st}]$;
- *the non-monotonic read-write pattern(RWP):*
  3) $r'' = R(w)$;
  4) $ver(w') \neq ver(w)$.

To calculate the probabilities of RI and WI, we view RI and WI as the concurrent occurrences of multiple atomicity violation patterns in the single-writer case. According to our previous work [2], atomicity violation in the single-writer case is equivalent to the pattern named Old New Inversion (ONI). We further destruct the ONI into the temporal Concurrency Pattern (CP) and the semantic Read Write Pattern (RWP). Then we obtain the probability of ONI:

$$
\begin{aligned}
\mathbb{P}\{\text{ONI}\} &= \sum_{m \geq 1} \mathbb{P}\{\text{CP} \mid R' = m\} \cdot \mathbb{P}\{\text{RWP} \mid R' = m\} \\
&\approx \sum_{m=1}^{n_r} \left(\left(\sum_{k=0}^{n_r-1} \binom{n_r}{k} \binom{m-1}{n_r-k-1} p_0^k r^{n_r-k} s^m \right) \right. \\
&\quad \cdot e^{-q\lambda_w t} \frac{\alpha^q B(q, \alpha(N-q)+1)}{B(q, N-q+1)} \\
&\quad \left. \cdot \left(1 - \left(\frac{J_1}{B(q, N-q+1)} \right)^m \right) \right).
\end{aligned}
\tag{2.2}
$$

where R' is a random variable denoting the number of r's in Definition 2.1. It is an approximation since the timed balls-into-bins model used for calculating the probability of read-write patterns assumes that there is at most one such r' in each reader queue, which can partly be justified by numerical results. For further details, please refer to our previous work [2] for detailed explanations of Equation 2.2.

Through observation we can see that ONI is a special case of RI when there exists single writer. Based on the quantification of ONI, we can quantify RI for MWMR registers with extra consideration of the concurrency effects of multiple writers. As for WI, by replacing r, w', r' in Definition 2.1 with w', w, r'' accordingly, it can be analyzed with similar process. Above all, we simplify MWMR quantification by reusing results of ONI and taking the extra effects of multiple writers into consideration together.

We first model the occurrences of multiple writers with the following queuing model. Each client's workload is recognized as an independent queue characterized by the rate of operations and the service time of each operation (i.e. $[o_{st}, o_{ft}]$). Assume a Poisson process with parameter λ for the scenario of each client issuing a sequence of read/write operations, and an exponential distribution with parameter

μ for the service time of each operation. We then have n independent, parallel $M/M/1/1/\infty/FCFS$ queues (i.e., a single-server exponential queuing system, whose capacity is 1 with the first come first served discipline) [3]. All the queues have arrival rate λ and service rate μ. For simplicity, queues of writes and reads are separate; and if there is any operation in service, no more operations can enter it in that queue.

Let $X^i(t)$ be the number of operations in queue i at time t. Then $X^i(t)$ is a continuous-time Markov chain with two states: 0 when the queue is empty and 1 when some operation is being served. Its stationary distribution $P_s \triangleq P(X^i(\infty) = s), s \in \{0, 1\}$ is:

$$
P_0 = \frac{\mu}{\mu + \lambda}, P_1 = \frac{\lambda}{\mu + \lambda}.
$$

Now we consider the concurrency conditions of RI and WI respectively. For RI, given a read r in Q_i, let W'_{cr} be a random variable denoting the number of w's satisfying $r_{st} \in [w'_{st}, w'_{ft}]$ in RI. The probability of the event that r starts during the service period of some write w' in Q_j equals the probability that when r arrives Q_i, it finds Q_i empty, as well as finds Q_j full as a bystander (with the constraint that Q_j is a writer queue). Since the events in different queues are independent, by the PASTA property and through combinatorial analysis, we have

$$
\mathbb{P}\{W'_{cr} = x\} = \binom{n_w}{x} \left(\frac{\lambda}{\mu + \lambda} \right)^x \left(\frac{\mu}{\mu + \lambda} \right)^{n_w - x + 1}
\tag{2.3}
$$

Conditioning on $W'_{cr} = w$, the probability of RI is no more than the sum of each w' forming RI with r seperately. Therefore, a probabilistic and combinatorial analysis shows that

$$
\begin{aligned}
\mathbb{P}\{\text{RI}\} &= \sum_{x=1}^{n_w} \mathbb{P}\{W'_{cr} = x\} \cdot \mathbb{P}\{\text{RI} \mid W'_{cr} = x\} \\
&\leq \sum_{x=1}^{n_w} \mathbb{P}\{W'_{cr} = x\} \cdot x \mathbb{P}\{\text{ONI}\}
\end{aligned}
\tag{2.4}
$$

As for WI, given a specific write w' in Definition 2.2, let $W_{cw'}$ be a random variable denoting the number of writes w satisfying $w'_{st} \in [w_{st}, w_{ft}]$ in RI. Similarly, through probabilistic and combinatorial analysis we have

$$
\mathbb{P}\{W_{cw'} = x\} = \binom{n_w - 1}{x} \left(\frac{\lambda}{\mu + \lambda} \right)^x \left(\frac{\mu}{\mu + \lambda} \right)^{n_w - x}
\tag{2.5}
$$

$$
\begin{aligned}
\mathbb{P}\{\text{WI}\} &= \sum_{x=1}^{n_w - 1} \mathbb{P}\{W_{cw'} = x\} \cdot \mathbb{P}\{\text{WI} \mid W_{cw'} = x\} \\
&\leq \sum_{x=1}^{n_w - 1} \mathbb{P}\{W_{cw'} = x\} \cdot x \mathbb{P}\{\text{ONI}\}
\end{aligned}
\tag{2.6}
$$

Substituting Formula (2.2)-(2.6) into (2.1), we obtain an upper bound of the rate of violating atomicity incurred in W2R1.

$\mathbb{P}\{\text{Atomicity Violation in W2R1}\}$

$$\leq \sum_{x=1}^{n_w} \mathbb{P}\{W'_{cr} = x\} \cdot x \mathbb{P}\{\text{ONI}\} + \sum_{x=1}^{n_w-1} \mathbb{P}\{W_{cw'} = x\} \cdot x \mathbb{P}\{\text{ONI}\}$$

$$\approx \frac{(2n_w - 1)\lambda\mu}{(\lambda+\mu)^2} \sum_{m=1}^{n_r} \left(\left(\sum_{k=0}^{n_r-1} \binom{n_r}{k} \binom{m-1}{n_r-k-1} p_0^k r^{n_r-k} s^m \right) \right.$$

$$\left. e^{-q\lambda_w t} \frac{\alpha^q B(q, \alpha(N-q)+1)}{B(q, N-q+1)} \left(1 - \left(\frac{J_1}{B(q, N-q+1)} \right)^m \right) \right)$$

The results in Fig. 7 are obtained when setting $\lambda = \mu = 10s^{-1}$ and $\lambda_r = \lambda_w = 20s^{-1}$. The source code for calculation of the numerical results can be found in our open source project on line [2]. The numerical results show that the probability of atomicity violation is quite low (mostly below 0.1% and can be below 10^{-9}) and will decrease when the number of replicas increases. Moreover, the probability of the concurrency pattern is close to 1. The probability of the read-write pattern is significantly less and decreases as the number of replicas increases. The probability of the read-write pattern ensure that the probability of the atomicity violation is almost zero. Besides, the probability of atomicity violation using the W2R1 algorithm has positive correlation to the number of concurrent writers. The comprehensive experimental results in Section 4 further confirm our theoretical analysis.


Fig. 7. The probabilities of atomicity violation.

3 W1R2 AND W1R1: THEORETICAL ANALYSIS AND EXPERIMENTAL EVALUATION

From the diamond schema we have one-round-trip write algorithms W1R2 & W1R1, as shown in Algorithm 5 & 6, which complete each write in only one communication round-trip. When a writer client initiates a write request, it would constructs a version for the written key-value pair based on its local sequence number. The version assigned for its written value is only guaranteed to be larger than any version returned by its previous reads or assigned by its previous writes, but not others. Therefore, one-round-trip write works well only for *single-writer client* registers since the only writer client can promise to preserve the version for a register monotonically increasing without conflicting with other writer clients. For *multiple-writer client* registers,

2. https://github.com/Lingzhi-Ouyang/Almost-Strong-Consistency-Cassandra/tree/master/numerical_results

one-round-trip write may result in *write-version inversion anomalies*: a later write updates a value with a smaller version than a previous write from a different writer client. Here, we define a write is *invisible* if the assigned version of the write is smaller than a quorum of replicas before the invocation of the write. The value updated by an invisible write cannot be return by any read.

In this section, we tend to demonstrate that one-round-trip write algorithms are useless from two aspects. First we prove both W1R2 and W1R1 don't guarantee the value returned by each read can be within bounded staleness in Section 3.1. Then we argue that both algorithms lead to frequent unsuccessful updates by quantifying the rate of invisible writes in Section 3.2.

Algorithm 5: W1R2

1 Code for client process $p_i (0 \le i \le n-1)$:

2 **procedure** OneRoundWRITE($key, value$)
3    $localSeq[key] \leftarrow localSeq[key]+1$
4    $version \leftarrow (localSeq[key], i)$
5   update($key, value, version$)

6 **procedure** TwoRoundREAD(key)
7    $replicas \leftarrow$ query(key)
8    $version \leftarrow$ maxVer($replicas$)
9    $value \leftarrow$ valWithMaxVer($replicas, version$)
10    $localSeq[key] \leftarrow version.seq$ ▷ Recommended.
11   update($key, value, version$)
12   **return** $value$

Algorithm 6: W1R1

1 **procedure** OneRoundWRITE($key, value$)
2    $localSeq[key] \leftarrow localSeq[key]+1$
3    $version \leftarrow (localSeq[key], i)$
4   update($key, value, version$)

5 **procedure** OneRoundREAD(key)
6    $replicas \leftarrow$ query(key)
7    $version \leftarrow$ maxVer($replicas$)
8    $value \leftarrow$ valWithMaxVer($replicas, version$)
9    $localSeq[key] \leftarrow version.seq$ ▷ Recommended.
10   **return** $value$

3.1 Proof of Atomicity Violation in One-Round-Trip Write

We prove that, in the W1R2 or W1R1 algorithm, the value returned by each read can be any stale. That is,

Theorem 3.1. *For MWMR registers, given* $k \in \mathbb{N}^*$, *the W1R2 or W1R1 algorithm doesn't satisfy k-atomicity.*

Proof. Assume W1R2 or W1R1 satisfies k-atomicity, where $k \in \mathbb{N}^*$. Then, we can raise a counterexample to show the absurdity of above assumption. The counterexample is constructed as follows. Suppose there exists two writer clients whose identifiers are p_0 and p_1, and we only focus on the same register. At time t_0, p_0 stores v_0 as its local sequence number for the register, and p_1 stores v_1, satisfying

$v_1 > v_0$. During the period from t_0 to t_1, p_1 executes k writes successively while p_0 initiates no writes. Thus, at time t_1, p_1's local version is increased to be $(v_1. + k, 1)$ while p_0 still keeps v_0. After that, during the period from t_1 to t_2, only p_0 executes k writes successively and its local version is increased to be $v_0 + k$ at t_2. Since $(v_0 + k, p_0) < (v_1 + k, p_1)$, the written value versioned $(v_0 + k, p_0)$ is invisible. Then at t_3, a read occurs but still reads the value versioned $(v_1 + k, p_1)$. In above case, the permutation of operations on p_0 and p_1 goes like this: (The *version* that a *write* assigns or a *read* returns are displayed in brackets.) $w(v_1 + 1, p_1) \prec_\pi$ $w(v_1 + 2, p_1) \prec_\pi ... \prec_\pi \underline{w(v_1 + k, p_1)} \prec_\pi w(v_0 + 1, p_0) \prec_\pi$ $w(v_0 + 2, p_0) \prec_\pi ... \prec_\pi \underline{w(v_0 + k, p_0)} \prec_\pi \underline{r(v_1 + k, p_1)}$. Since $r(v_1 + k, p_1)$ returns a value not written by one of the latest k preceding writes in π, so the execution violates k-atomicity. By contradiction, we prove the correctness of Theorem 3.1. □

3.2 Quantification of Visible Writes in One-Round-Trip Write

W1R2 & W1R1 not only don't satisfy bounded k-atomicity, but also lead to frequent invisible writes that cannot be read. Here, we quantify the rate of visible writes incurred in W1R2 & W1R1 using the following assumption. Assume a Poisson process with parameter λ for the scenario of each client issuing a sequence of write/read operations, but *ignore* the duration of each operation. That is, each write/read is regarded as an instant event taking effect on the register at its starting moment. Suppose there are n_w writer clients($n_w > 1$), identified 0, 1, ..., $n_w - 1$ separately. Then we have n_w independent, parallel write sequences, each satisfying a Poisson process with parameter λ. Assume that at the initial moment, all n_w writer clients are informed of the version (v_0, id) for some specific register that we focus on here. Let $N^i(t)$ be the number of write operations issued by writer client i from the initial moment till time t, where $0 \le i \le n_w - 1$. According to the formula of Poisson probability, we have

$$\mathbb{P}\{N^i(t) = k\} = \frac{(\lambda t)^k e^{-\lambda t}}{k!}$$

Assume that writer client i stores the local sequence number $k - 1$ at time t. Let $V^i(t, k)$ be a random variable denoting the visibility of the un-occurred write versioned $(v_0 + k, i)$ at time t. Then, $V^i(t, k)$ has two status: invisible(denoted 0) and visible(denoted 1). $V^i(t, k) = 1$ occurs if and only if all other writer clients' local sequence number stayed smaller than $(v_0 + k, i)$ at time t. Since write sequences are independent, we have

$$\mathbb{P}\{V^i(t, k) = 1\}$$
$$= \prod_{j=0}^{i-1} \mathbb{P}\{N^j(t) < k+1\} \prod_{j=i+1}^{n_w-1} \mathbb{P}\{N^j(t) < k\}$$
$$= \left(\sum_{j=0}^{k} \frac{(\lambda t)^j e^{-\lambda t}}{j!} \right)^i \left(\sum_{j=0}^{k-1} \frac{(\lambda t)^j e^{-\lambda t}}{j!} \right)^{n_w-1-i}$$

Here, we quantify the rate when $k = 1$ at the expected arrival time $t = \frac{1}{\lambda}$, when the probability of $N^j(t) = 1$ is the highest.

By setting $\lambda = 10s^{-1}$, we present the numerical results in Fig. 8 and Table 3. The numerical results show that the probability of visible writes will fall as there exist multiple writer clients. When the number of writer clients is up to 10, the probability of visible writes for the writer client with the maximum id is about 0.063, while for the writer client with the minimum id, the probability is about 10^{-4}.

Fig. 8. The probabilities of visible writes from writer clients with different identifiers ($\lambda = 10s^{-1}, t = 100ms$).
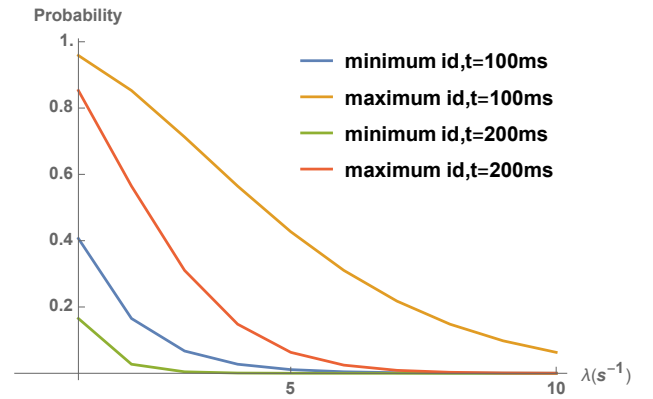

Fig. 9. The probabilities of visible writes from the writer client with maximum / minimum identifier by varying λ ($n_w = 10$).

3.3 Experimental Results of One-Round-Trip Write Algorithms

The rough theoretical analysis in Section 3.2 doesn't take the duration of each operation into consideration, so we conduct more comprehensive experimental evaluations of W1R2 and W1R1 in quorum-replicated datastores. The experimental design and the environment settings are described in the experiments of W2R1, and the experimental results of W1R2 and W1R1 are presented in Table 4.

For read operations, the overall probability of obtaining an up-to-date value is varied from 34.4% to 91.1%. However, under most environment settings, the probability is around 50%. That is, nearly half read operations will return some stale values. Moreover, the staleness of the returned data can be quite high. We only calculate the rate of k up to 10 in terms to k-atomicity. Specifically, the probability of returning an value with the staleness of 5 is varied from 0.7% to 8.1%. When the staleness is up to 10, the probability

TABLE 3
Numerical results on the probabilities of **invisible** writes in one-round-trip write algorithms without extra information in acknowledgement($\lambda = 10s^{-1}, t = 100ms$).

Writer numbers	$id=0$	$id=1$	$id=2$	$id=3$	$id=4$	$id=5$	$id=6$	$id=7$	$id=8$	$id=9$
$n_w=2$	0.632121	0.264241								
$n_w=3$	0.864665	0.729329	0.458659							
$n_w=4$	0.950213	0.900426	0.800852	0.601703						
$n_w=5$	0.981684	0.963369	0.926737	0.853475	0.70695					
$n_w=6$	0.993262	0.986524	0.973048	0.946096	0.892193	0.784386				
$n_w=7$	0.997521	0.995042	0.990085	0.98017	0.96034	0.92068	0.84136			
$n_w=8$	0.999088	0.998176	0.996352	0.992705	0.98541	0.97082	0.94164	0.883279		
$n_w=9$	0.999665	0.999329	0.998658	0.997316	0.994633	0.989265	0.97853	0.957061	0.914122	
$n_w=10$	0.999877	0.999753	0.999506	0.999013	0.998025	0.996051	0.992102	0.984204	0.968407	0.936814

TABLE 4
Experimental results of one-round-trip write algorithms without extra information in acknowledgement($\lambda = 10s^{-1}, t = 100ms$).

Algorithm	Quorum level	Replica	Inter-DC delay/ms	Client num.	Read ratio	$\mathbb{P}(k=1)$	$\mathbb{P}(k=5)$	$\mathbb{P}(k=10)$
W1R2	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	10	0.90	0.344	0.081	0.017
W1R2	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	20	0.90	0.452	0.084	0.010
W1R2	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	30	0.90	0.541	0.058	0.011
W1R2	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	40	0.90	0.624	0.057	0.020
W1R2	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	50	0.90	0.676	0.050	0.007
W1R2	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	30	0.50	0.911	0.007	0.010
W1R2	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	30	0.60	0.878	0.014	0.009
W1R2	QUORUM	3	$\mathbb{N}(50, 25^2)$	30	0.90	0.494	0.057	0.006
W1R2	QUORUM	3_1_1	$\mathbb{N}(50, 25^2)$	30	0.90	0.518	0.069	0.004
W1R2	QUORUM	3_3_3	$\mathbb{N}(50, 25^2)$	30	0.90	0.529	0.064	0.016
W1R2	EACH_QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	30	0.90	0.540	0.071	0.016
W1R1	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	10	0.90	0.375	0.083	0.010
W1R1	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	20	0.90	0.485	0.085	0.013
W1R1	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	30	0.90	0.507	0.052	0.023
W1R1	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	40	0.90	0.591	0.057	0.021
W1R1	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	50	0.90	0.618	0.077	0.026
W1R1	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	30	0.50	0.864	0.017	0.012
W1R1	QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	30	0.60	0.816	0.024	0.021
W1R1	QUORUM	3	$\mathbb{N}(50, 25^2)$	30	0.90	0.496	0.091	0.009
W1R1	QUORUM	3_1_1	$\mathbb{N}(50, 25^2)$	30	0.90	0.466	0.061	0.011
W1R1	QUORUM	3_3_3	$\mathbb{N}(50, 25^2)$	30	0.90	0.563	0.059	0.014
W1R1	EACH_QUORUM	1_1_1	$\mathbb{N}(50, 25^2)$	30	0.90	0.496	0.053	0.019

TABLE 5
Numerical results on the probabilities(lower bound) of **invisible** writes in one-round-trip write algorithms if sequence number is included in acknowledgement($\lambda = 10s^{-1}, t = 100ms$).

Writer numbers	$id=0$	$id=1$	$id=2$	$id=3$	$id=4$	$id=5$	$id=6$	$id=7$	$id=8$	$id=9$
$n_w=2$	0.264241	0.0803014								
$n_w=3$	0.458659	0.323324	0.154154							
$n_w=4$	0.601703	0.502129	0.377662	0.222077						
$n_w=5$	0.70695	0.633687	0.542109,	0.427636	0.284545					
$n_w=6$	0.784386	0.730482	0.663103	0.578878	0.473598	0.341997				
$n_w=7$	0.84136	0.8017	0.752125	0.690156	0.612695	0.515869	0.394836			
$n_w=8$	0.883279	0.854099	0.817624	0.77203	0.715037	0.643796	0.554745	0.443431		
$n_w=9$	0.914122	0.892652	0.865815	0.832269	0.790336	0.73792	0.6724	0.5905	0.488125	
$n_w=10$	0.936814	0.921018	0.901272	0.87659	0.845738	0.807172	0.758965	0.698707	0.623383	0.529229

is around 1%. Note that if the returned data is with staleness of 5 or even more, it is nearly of no use for users. Thus, reader clients may frequently obtain stale or even invalid data. Besides, for write operations, there's still a considerable possibility to issue invisible values. Therefore, the W1R2 and W1R1 algorithms without any improvement are far away from use. [3]

3.4 Improvement of One-Round-Trip Write Algorithms

Imagine slightly modifying the procedure of $update$ function. When a replica receives an "update" message and finds

3. For MWMR registers, the W1R2 and W1R1 algorithms that use discrete version (seq, id) are far away from use. However, the one-round-trip write algorithms in Cassandra and some other datastores use timestamps following the UTC standard, which is different from our algorithms. Theoretically speaking, these datastores cannot always promise to produce unique timestamps for different data, and that's out of scope in this work.

the given version smaller than it stores, then it would fill the acknowledgement with extra information, like the sequence number of the stored version and reply to the client. A client would have to wait for acknowledgement from a majority of replicas and update its local sequence number according to the information in acknowledgement. Then, an invisible one-round-trip write may work like the $query$ function and help update the local sequence number on the client side. In this way, W1R2 can be proved to satisfy n_w-atomicity given n_w writer clients.

Here we quantify the rate of visible writes in the improved W1R2 & W1R1 algorithms. Suppose the writer client i issues a write w after time t from its last write operation w'. Let $U^i(t)$ be a random variable denoting the visibility of the write w. Similarly, $U^i(t)$ has two status: invisible(denoted 0) and visible(denoted 1). Here, we only focus on a specific but frequent pattern that would inevitably make w invisible: during the period from the finish time of w' to the invocation of w, there exists a writer client identified smaller than i completing at least three writes, or a writer client identified larger than i completing at least two writes [4]. Since write sequences are independent, through probabilistic and combinatorial analysis we have

$$\mathbb{P}\{U^i(t) = 0\}$$
$$\geq 1 - \prod_{j=0}^{i-1} \mathbb{P}\{N(t) < 3\} \prod_{j=i+1}^{n_w-1} \mathbb{P}\{N(t) < 2\}$$
$$= 1 - e^{-\lambda t(n-1)} \cdot (1 + \lambda t + \frac{1}{2}(\lambda t)^2)^i \cdot (1 + \lambda t)^{n_w-1-i}$$

Then, we have

$$\mathbb{P}\{U^i(t) = 1\} = 1 - \mathbb{P}\{U^i(t) = 0\}$$
$$\leq e^{-\lambda t(n-1)} \cdot (1 + \lambda t + \frac{1}{2}(\lambda t)^2)^i \cdot (1 + \lambda t)^{n_w-1-i}$$

Here, we quantify the rate at the expected time $t = \frac{1}{\lambda}$, which has the highest probability of that there exists no other operations between w' and w.

By setting $\lambda = 10s^{-1}$, we present the numerical results in Fig. 10 and Table 5. The numerical results of the improved W1R2 & W1R1 algorithms are similar to those without improvement in Section 3.2, although the probability of visible writes raises a little. When the number of writer clients is up to 10, the probability of visible writes for the writer client with the maximum id is about 0.047, while for the writer client with the minimum id, the probability is about 0.063. However, the improvement is not significant. Actually, the modification in *update* doesn't prevent the occurrence of invisible writes, so the problem of the one-round-trip write algorithm still exists.

4 W2R1: EXPERIMENTAL EVALUATIONS

4.1 ASC Tradeoff in Different Environments

We investigate the impacts of different environments on the performance of the W2R1 algorithms running on Cassandra. The W2R2 algorithm is also conducted for comparison. The environmental settings are varied from workload pattern (client number and read ratio), replica configuration, communication delay to read quorum level.

4. This can be proved through a case-by-case analysis.

Fig. 10. The probabilities of visible writes from writer clients with different identifiers using the improved update acknowledgement ($\lambda = 10s^{-1}, t = 100ms$).


Fig. 11. The probabilities of visible writes from the writer client with maximum / minimum identifier by varying λ using the improved update acknowledgement ($n_w = 10$).

4.1.1 Impact of Client Number

The client number reflects the concurrency patterns of operations. Intuitively, more clients will result in higher concurrency. In the experiment, each client is acted by one YCSB instance. All clients are allowed to initiate read/write requests concurrently, while each of them can only execute at most one request at any time. By varying client number from 10 to 40, we derive the results shown in Fig. 12.

First we have a look at the consistency performance of each algorithm. W2R2 guarantees atomicity all the time as is expected, while other algorithms can lead to atomicity violation more or less. From the perspective of worst case we produced, W2R1 produced some non-atomic traces with k up to 4, while others results in a better performance. From the perspective of average consistency, W2R1 also has the worst performance compared with other three algorithms that applies Cassandra's mechanisms(digest request, read repair or snitch strategy). However, although all but W2R2 can't satisfy atomicity all the time, they guarantee atomicity most of the time, with a confidence that more than 99.97% read requests can obtain the most up-to-date value. In overall, the average consistency performance become weaker as the concurrent client number increases, and we produced the worst case of $k = 4$ when client number is up to 40. We speculate that the rate of returning fresh data for reads would decrease with more concurrent clients operating on single register, while all above non-atomic algorithms can also promise atomicity most of the time when concurrent

clients are no more than 50, which is applicable for many applications and scenarios.

As for latency, we observe that average request latency is positively relative to client number in a subtle degree. The main reason is obvious: the processing capacity of each server is limited and would be gradually exhausted as the user throughput increases. From another aspect, W2R2 cost higher *read* latency compared with other four algorithms, while it gains stronger consistency all the time. Different algorithms behaves in a slightly different way and we'll discuss it in detail in Section 4.2.

Fig. 12. Consistency and latency results by tuning the number of clients.

4.1.2 Impact of Read Ratio

In the experiments, each client can initiate read and write requests. The ratio between read and write can also be a critical factor to impact the chance to produce violation patterns. Here, we use *read proportion*(the ratio of read numbers to total operation numbers) to describe the relations between read and write numbers. By varying read proportion from 0.50 to 0.99, we derive the results shown in Fig. 13. As we can see, W2R2 still guarantees atomicity while others do not. W2R1 with snitch leads to a monotonically decreasing atomicity with higher read proportion. For W2R1, digest and repair algorithms, higher read proportion within certain range (0.5 - 0.9) can result in weaker consistency in average. However, when the read proportion is up to 0.99, the consistency performance reversely becomes better. One potential reason is that, k-atomicity violation are described by the staleness of values returned by reads, so within certain range, more reads will have more chances to form more complicated read-write concurrent patterns and return stale results. However, as the ratio of reads grows too high, the write will become too rare and sparsely distributed for reads to obtain stale written values.

As for latency, we observe that average request latency is approximately irrelative to read proportion for all algorithms except W2R1 with digests. The reason why request latency for W2R1 with digests grows higher when read proportion gets lower is mainly that, when write proportion becomes relatively higher, more frequent updates will make digest mismatch occur more frequently, triggering extra round-trip for collecting full data. If we only focus on *read* latency, W2R2 cost the highest *read* latency among all, while it guarantees strong consistency. Different algorithms behaves in a slightly different way and we'll discuss it in detail in Section 4.2.

4.1.3 Impact of Replica configuration

The replica factor in Cassandra specifies the number of replicas in each data center. The number of replicas impacts


Fig. 13. Consistency and latency results by tuning read ratio.

the consistency and latency performance in some degree, while the data center location of replicas impacts the request latency. Given a number of replicas, besides spreading them into different data centers, we can also place them all into a single data center for the convenience of local access. In the experiments, the replica factor is varied in 3(inside single data-center), 1_1_1, 3_1_1, and 3_3_3. As is shown in Fig. 14, more replicas would result in stronger consistency and higher latency in overall. The reason is obvious: the latency of an operation depends on the slowest responding replica which costs the longest time to communicate. In the network where delays are randomly distributed, it's more probable to have longer delays when communicating with more replicas. However, W2R1 *with snitch* behaves in a unique way. Applying smart routing through the snitch strategy makes the algorithm always process each read/write request by accessing local replicas in priority, which tends to lower the overall latency, but replica updates in remote data centers are usually not able to spread in time. 3_1_1 is a powerful evidence for above explanation.


Fig. 14. Consistency and latency results by tuning replica configuration.

4.1.4 Impact of Network Delay

In our experiments, we mainly focus on inter-data center delay because it impacts the performance in a significant way compared to intra-data center delay. The default one-way inter-data center delays in network are normally distributed in the form of $\mathbb{N}(\mu, \sigma^2)$, with $\mu = 50$ and $\sigma = 25^2$. we explore the impacts of network delay by varying μ from $10ms$ to $50ms$, and σ from $5ms$ to $25ms$ respectively. From Fig. 15 and Fig. 16 we observe that, the average latency of write/read operations are basically in proportion to average network delays among servers. Besides, the consistency performance become weaker when average network delays or jitters grow higher and higher. This is due to, not only higher network delay would make each operation duration stay longer, but also larger jitters or variances of delays would aggravate the occurrence of out-of-sync replicas. Then, the

joint effects of above both lead to frequent atomicity violation patterns.


Fig. 15. Consistency and latency results by tuning the average value of inter-DC delay.

Fig. 16. Consistency and latency results by tuning inter-DC delay jitter.

4.1.5 Impact of Read Quorum Level

We test two types of read quorum level in Cassandra: QUORUM and EACH_QUORUM [5]. The level of QUORUM requires to access a quorum of replicas from all data-centers, which is actually the same as what we mention in the theoretical model. The level of EACH_QUORUM requires to access a quorum of replicas from each data-centers, which is actually the stronger than the requirement of theoretical model. As is shown in Fig. 17, EACH_QUORUM promises better consistency performance in average at the cost of extra latency; QUORUM lowers the average latency at the sacrifice of a little bit lower consistency rate. The tradeoff between consistency and latency is well presented in this experiment.


Fig. 17. Consistency and latency results by tuning read quorum level.

5. Cassandra also provides a special quorum level: LOCAL_QUORUM. The level of LOCAL_QUORUM only requires to access a quorum of replicas in the local data-center, which speeds up a lot while does actually not meet the requirement of majority-replication. Therefore, it's out of scope for this work.

4.2 ASC Tradeoff with Different Algorithms

From all above experimental results by varying different factors, we observe that distinct algorithms vary in performance. Here, we summarize some insights from the following several perspectives.

- **Consistency** No matter Under what circumstance, W2R2 guarantees atomicity all the time, while other algorithms do not. Although from the staleness of returned data, other algorithms produced non-atomic traces with k up to 4 at worst, the probability of that is super low. Specifically, all algorithms except W2R2 produce stale reads with the probability lower than 0.3% in our experiments under all circumstances, and lower than 0.02% under most circumstances. For example, from the experimental results under the default setting in Table 6, the rate of 2-atomic reads is lower than 0.03%, and 3-atomic reads even lower, not to mention higher values of k in terms of k-atomic reads. Specifically, W2R1 with repair performs best except W2R2, and its probability of atomic reads in average is merely a little inferior than W2R2. W2R1 with digests or snitch don't perform so well as W2R1 with repair, but better than W2R1 in overall. Thus, we believe that all these algorithms rarely violate atomicity, and still promise atomicity most of the time.

- **Latency** We mainly focus on read latency. Normally, W2R1 presents lower read latency than W2R2, which is expected. Surprisingly, W2R1 with digests costs the largest read latency in average. We speculate in two reasons. Firstly, using digest requests means that each time when processing read requests, only one replica would be read with full data, and all others with digests, which cannot promise to read in only one round-trip once there exists some inconsistent replicas. Secondly, compared to raw data collection, digest requests reduces data size in communication but requires extra processing work, which takes extra time. W2R1 with repair leads to approximately the same or just slightly higher read latency compared to W2R1 without repair. This is due to that it applies asynchronous mechanism to repair stale replicas in background after the response of read operations, which has nearly no effects on the read response time. As for W2R1 with snitch, it gains the highest efficiency for reads mainly because it always selects the replicas with high proximity and route requests to them, which significantly decreases latency.

Among all algorithms, W2R2 gains the strongest consistency at the cost of extra communication round-trip of each read, which inevitably leads to high read latency. By omitting the second round-trip of read, W2R1 without any optimization gains lower latency but the worst consistency performance in average, although it still guarantees atomicity most of the time.

Compared with the original W2R2 or W2R1 algorithms, the optimization in Cassandra can either help increase consistency through replica synchronization(*repair*), or speedup request processing by reducing communication data size(*digests*) or selecting replicas smartly(*snitch*). Applying these extra replica-related mechanisms in Cassandra

TABLE 6
ASC tradeoff under the default setting

	k_{max}	$\mathbb{P}(k=1)$	$\mathbb{P}(k=2)$	$\mathbb{P}(k=3)$	$\mathbb{P}(staleness)$	Read latency
W2R2	1	100%	0	0	0	100%
W2R1	3	99.9796%	0.0203%	0.0001%	0.0204%	53%
Snitch	2	99.9896%	0.0104%	0	0.0104%	52.3%
Digest	2	99.9926%	0.0074%	0	0.0074%	76.5%
Repair	2	99.9977%	0.0023%	0	0.0023%	53.9%

to W2R1 separately, we observe that W2R1 with snitch has the lowest read latency due to its smart routing, but produces weakest consistency among three. W2R1 with repair performs better in promising atomicity without increasing much latency in average since it puts repair work in the background. W2R1 with digests performs worse than W2R1 with repair but better than W2R1 with snitch in consistency, while its read latency is a lot higher than other one-round-trip read algorithms.

Among all, we think W2R1 with *repair* performs best concerning consistency and latency performance for two reasons. Firstly, the wisdom of this algorithm lies in the *divergence-oriented communication*. That is, only when inconsistency has been found will the algorithm take extra communication round-trip for repair, which costs less time for most consistent requests while still keeps a high degree of atomicity. Secondly, repairing replicas in background asynchronously is another core technique for optimizing both consistency and latency performance.

REFERENCES

[1] P. B. Gibbons and E. Korach, "Testing shared memories," *SIAM Journal on Computing*, vol. 26, no. 4, pp. 1208–1244, 1997.
[2] H. Wei, Y. Huang, and J. Lu, "Probabilistically-atomic 2-atomicity: Enabling almost strong consistency in distributed storage systems," *IEEE Trans. Comput.*, vol. 66, no. 3, pp. 502–514, Mar. 2017. [Online]. Available: https://doi.org/10.1109/TC.2016.2601322
[3] S. M. Ross, *Introduction to probability models*. Academic press, 2014.